

The No-Marketing Bullshit Introduction to Couchbase Server 2.0 and ext/couchbase

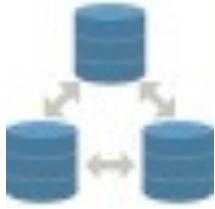
Jan Lehnardt
[@janl / jan@couchbase.com](mailto:@janl)

COUCHBASE



Couchbase Server Features

- Built-in clustering – All nodes equal



- Data replication with auto-failover

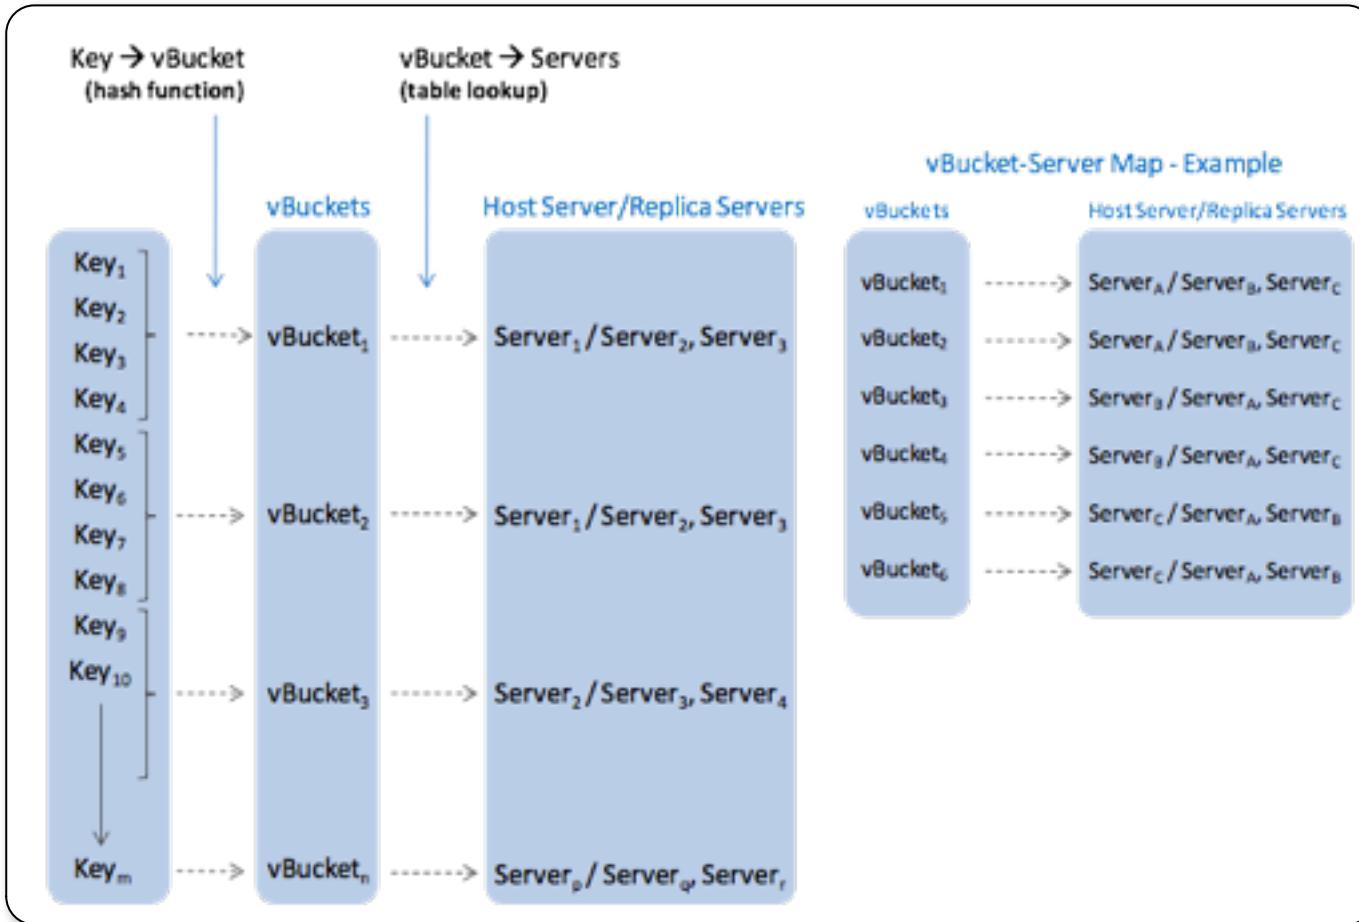


- Zero-downtime maintenance



- Clone to grow and scale horizontally

Auto-sharding: vBuckets



Couchbase Server Basic Operation

APP SERVER 1



APP SERVER 2



- Docs distributed evenly across servers in the cluster

SERVER 1



SERVER 2



SERVER 3



COUCHBASE SERVER CLUSTER

Couchbase Server Basic Operation

APP SERVER 1



APP SERVER 2



- Docs distributed evenly across servers in the cluster
- Each server stores both *active* & *replica* docs
 - Only one server active at a time

SERVER 1



SERVER 2



SERVER 3



COUCHBASE SERVER CLUSTER

Couchbase Server Basic Operation

APP SERVER 1



APP SERVER 2



- Docs distributed evenly across servers in the cluster
- Each server stores both *active* & *replica* docs
 - Only one server active at a time
- Client library provides app with simple interface to database

SERVER 1



SERVER 2

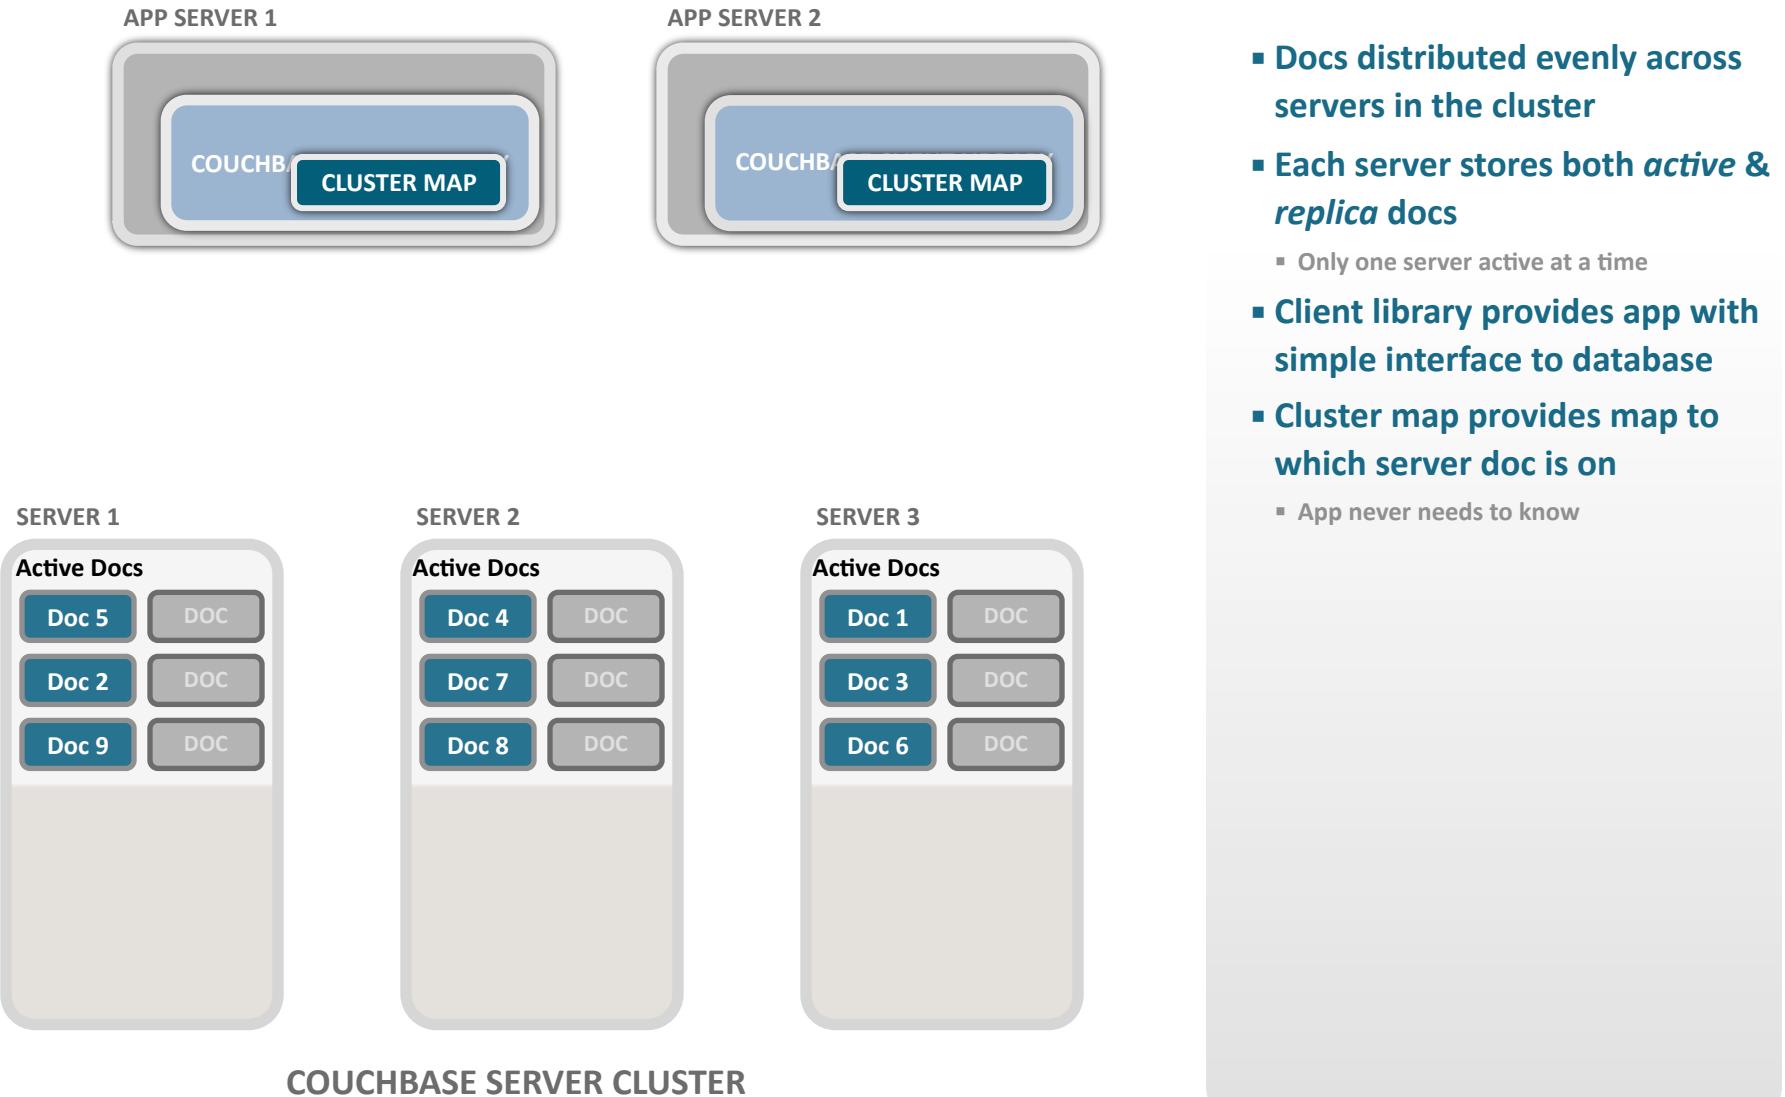


SERVER 3

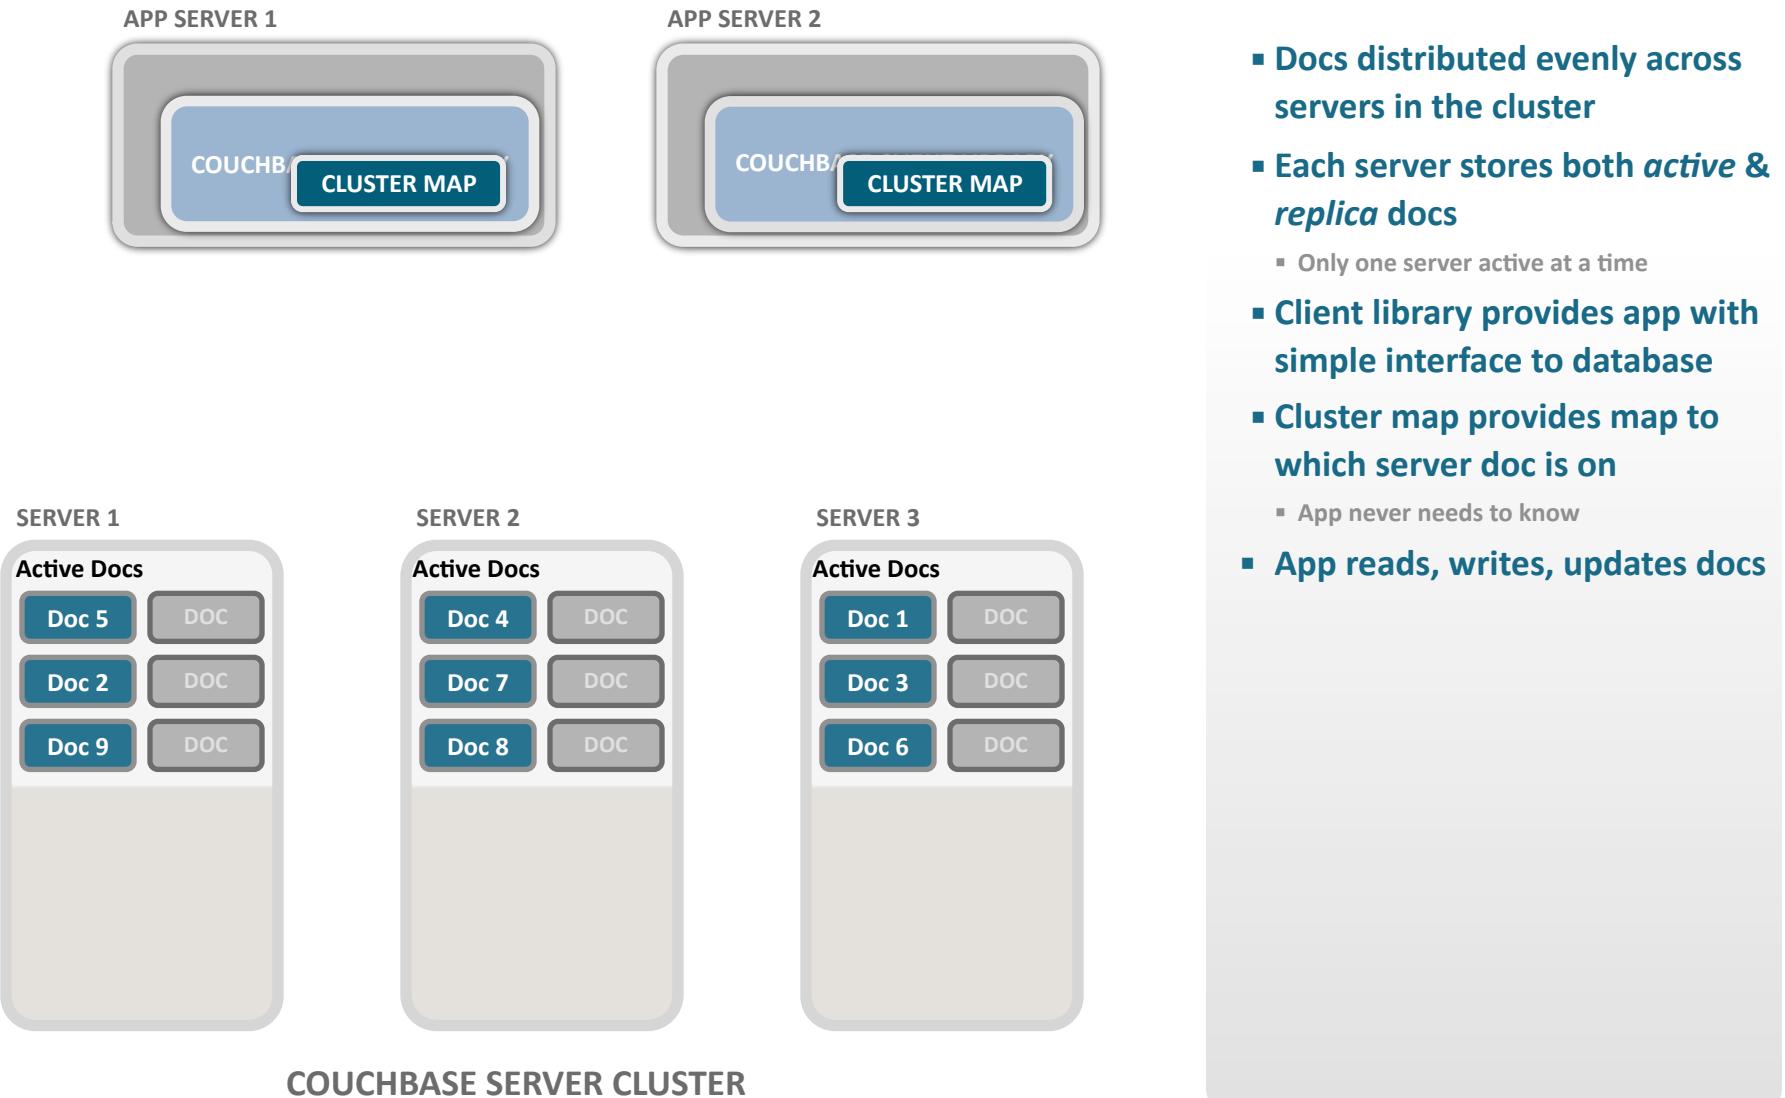


COUCHBASE SERVER CLUSTER

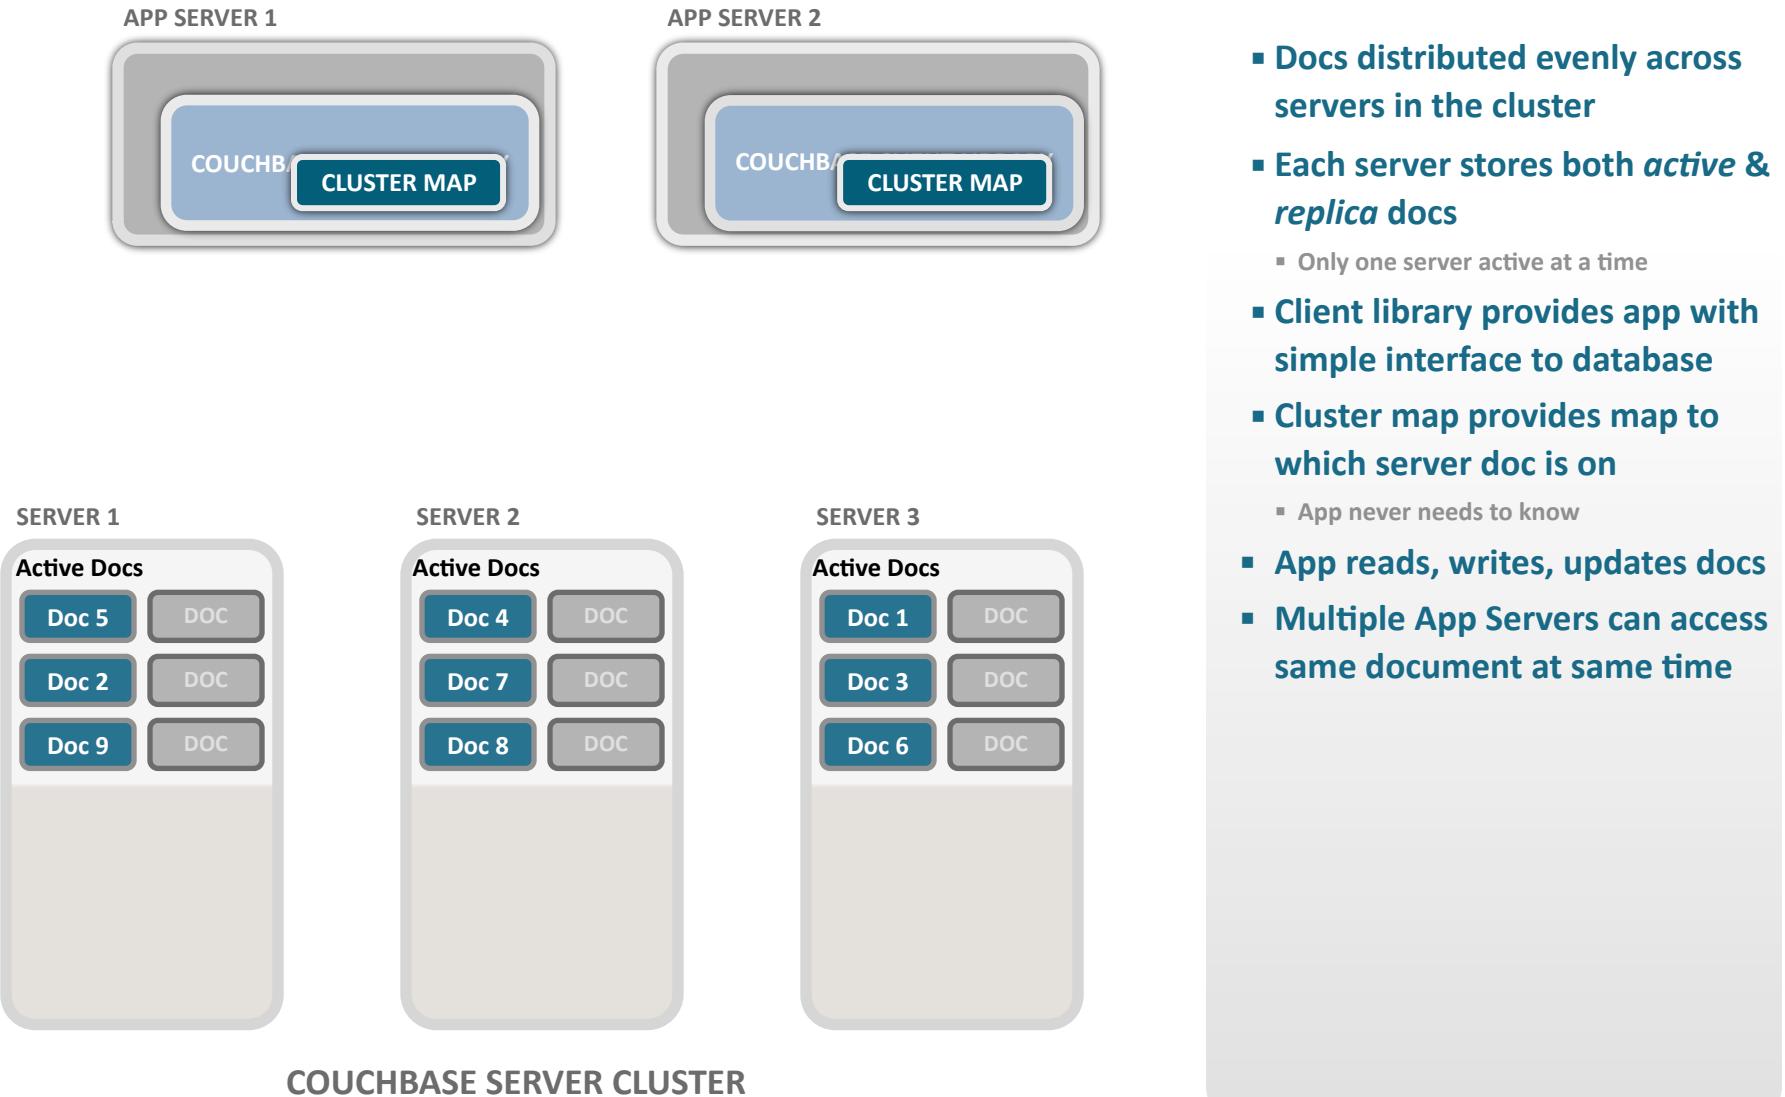
Couchbase Server Basic Operation



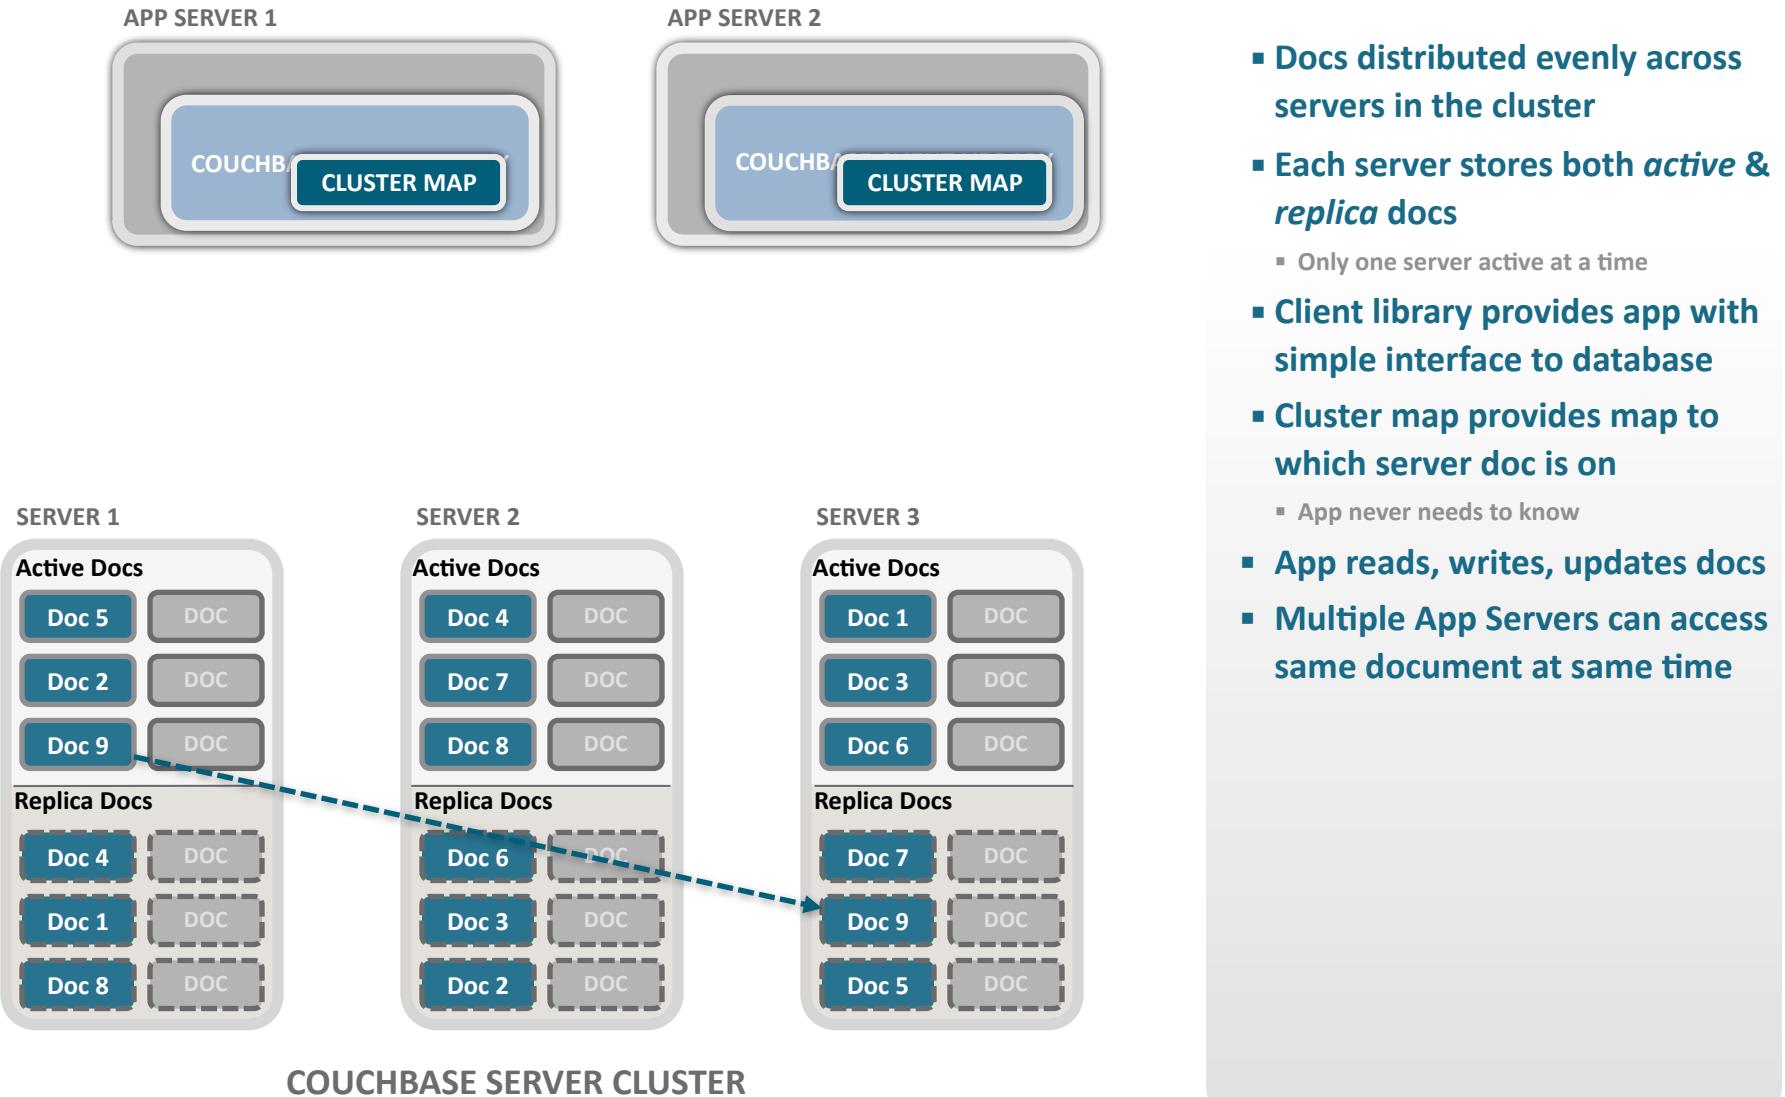
Couchbase Server Basic Operation



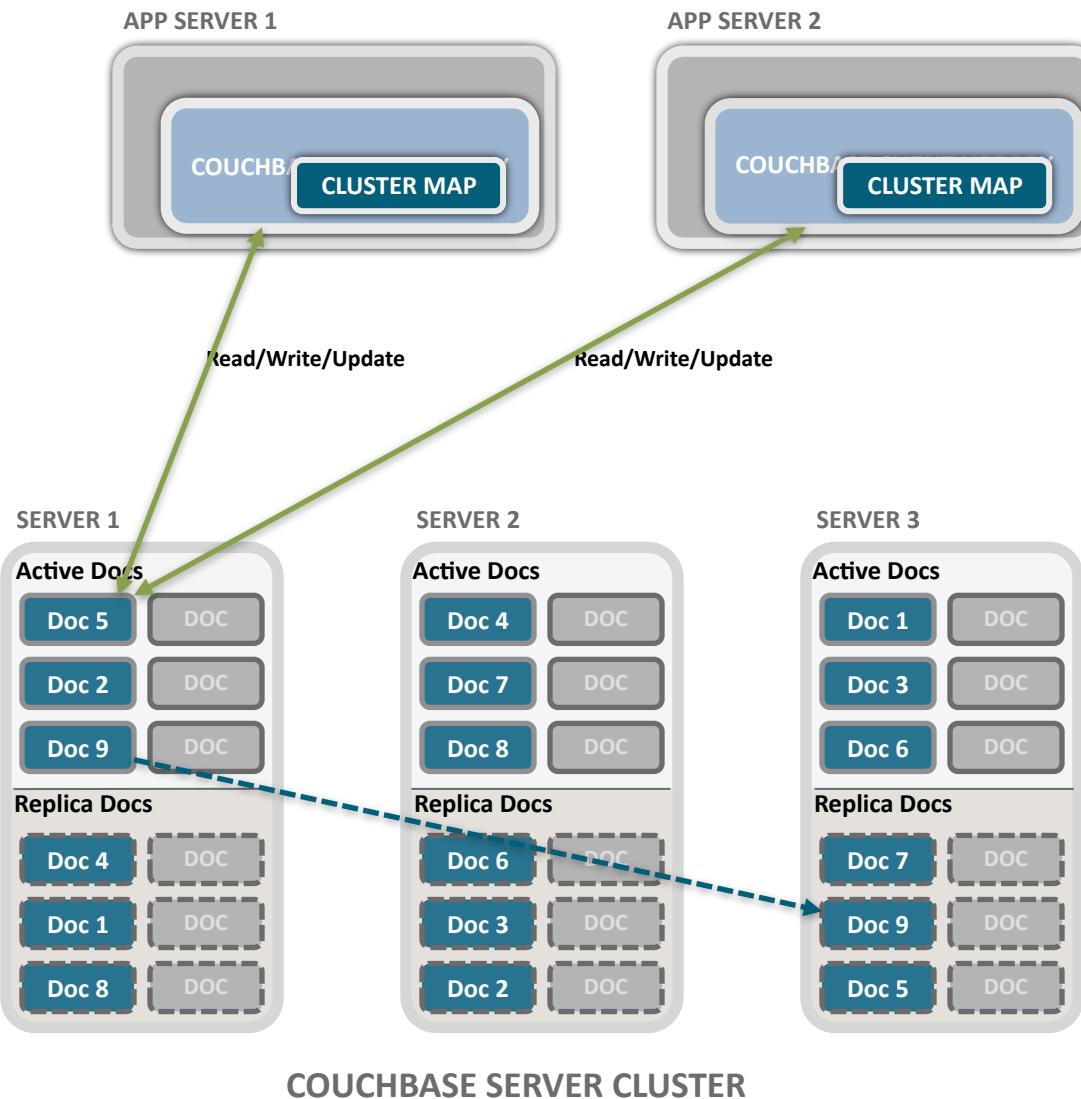
Couchbase Server Basic Operation



Couchbase Server Basic Operation

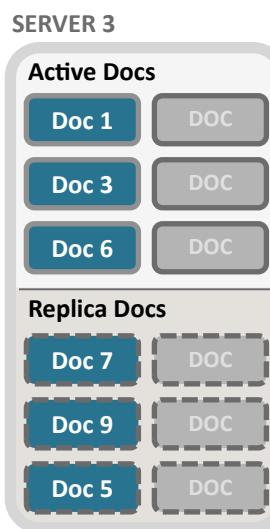
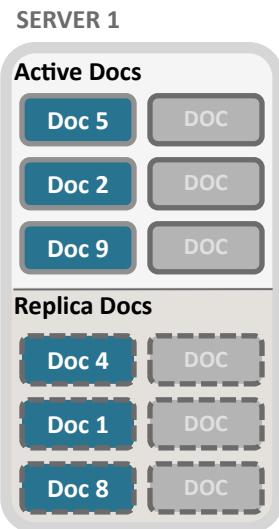
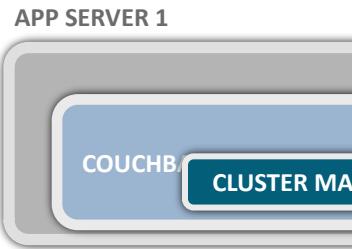


Couchbase Server Basic Operation



- Docs distributed evenly across servers in the cluster
- Each server stores both *active* & *replica* docs
 - Only one server active at a time
- Client library provides app with simple interface to database
- Cluster map provides map to which server doc is on
 - App never needs to know
- App reads, writes, updates docs
- Multiple App Servers can access same document at same time

Add Nodes



COUCHBASE SERVER CLUSTER

Add Nodes

APP SERVER 1



APP SERVER 2



SERVER 1



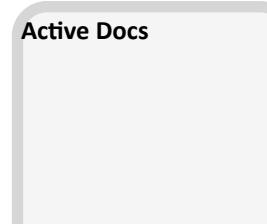
SERVER 2



SERVER 3



SERVER 4



SERVER 5



COUCHBASE SERVER CLUSTER

Add Nodes

APP SERVER 1



APP SERVER 2



- Two servers added to cluster
 - One-click operation

SERVER 1



SERVER 2



SERVER 3



SERVER 4



SERVER 5



COUCHBASE SERVER CLUSTER

Add Nodes

APP SERVER 1



APP SERVER 2



- Two servers added to cluster
 - One-click operation
- Docs automatically rebalanced across cluster
 - Even distribution of docs
 - Minimum doc movement

SERVER 1



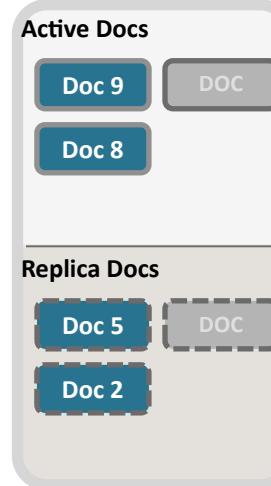
SERVER 2



SERVER 3



SERVER 4



SERVER 5



COUCHBASE SERVER CLUSTER

Add Nodes

APP SERVER 1



APP SERVER 2



- Two servers added to cluster
 - One-click operation
- Docs automatically rebalanced across cluster
 - Even distribution of docs
 - Minimum doc movement
- Cluster map updated

SERVER 1



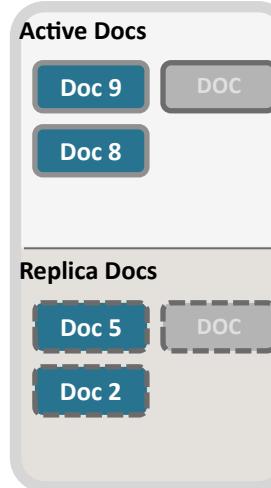
SERVER 2



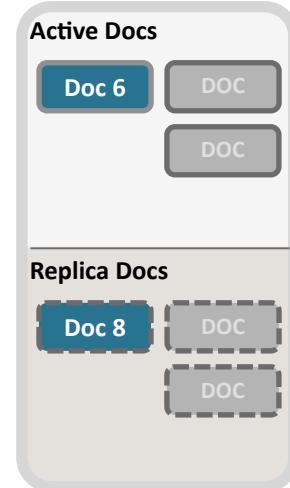
SERVER 3



SERVER 4



SERVER 5



COUCHBASE SERVER CLUSTER

Add Nodes

APP SERVER 1

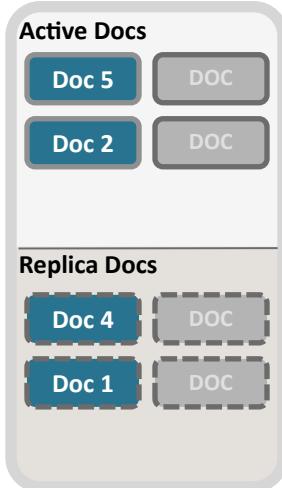


APP SERVER 2

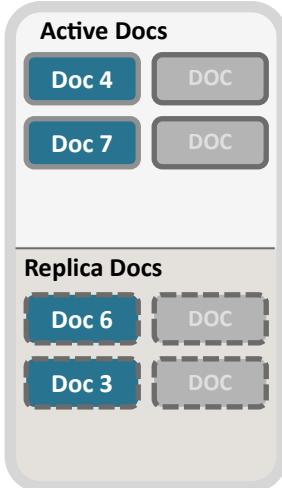


- Two servers added to cluster
 - One-click operation
- Docs automatically rebalanced across cluster
 - Even distribution of docs
 - Minimum doc movement
- Cluster map updated
- App database calls now distributed over larger # of servers

SERVER 1



SERVER 2



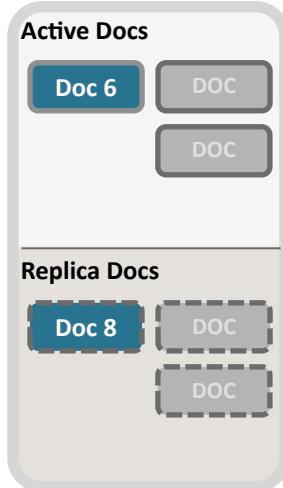
SERVER 3



SERVER 4

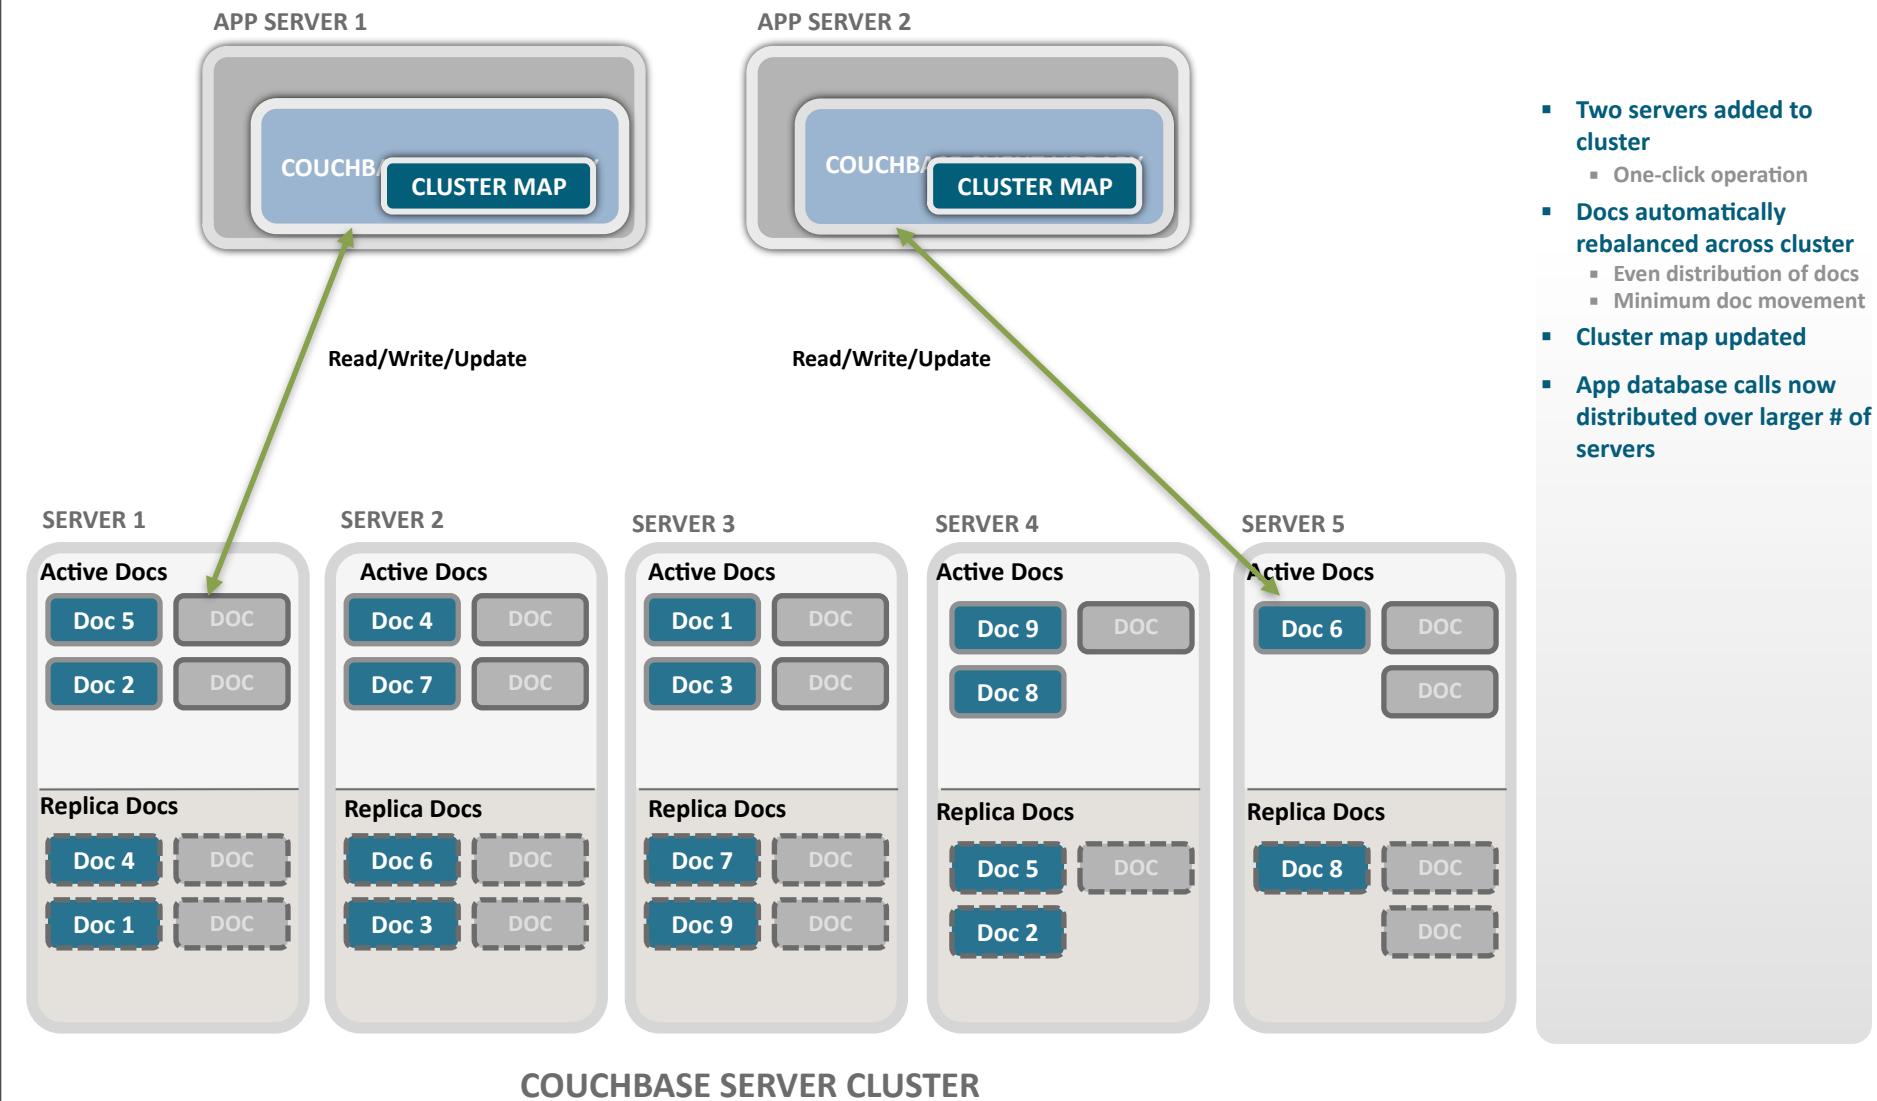


SERVER 5



COUCHBASE SERVER CLUSTER

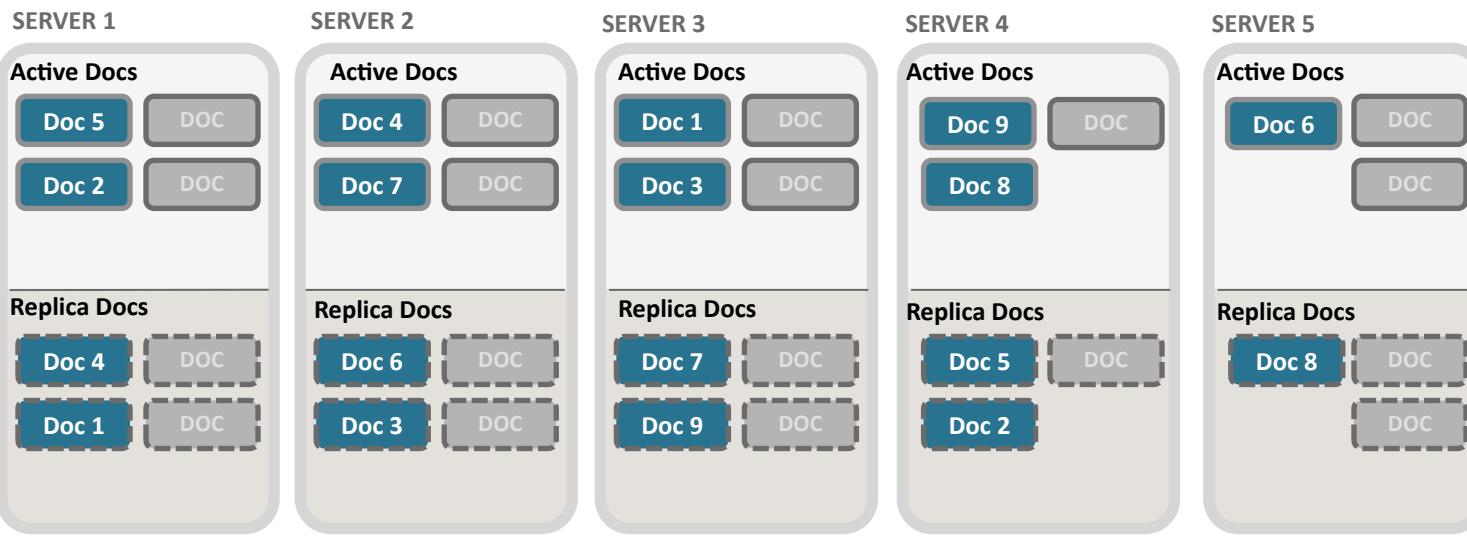
Add Nodes



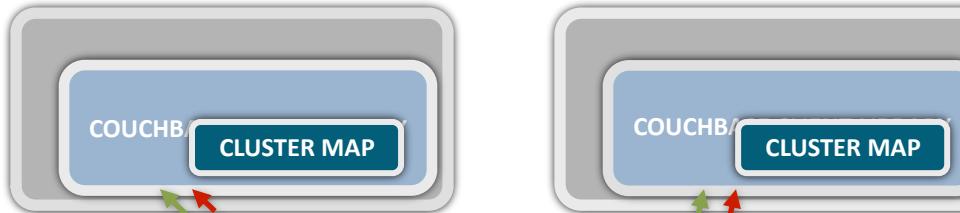
Fail Over Node



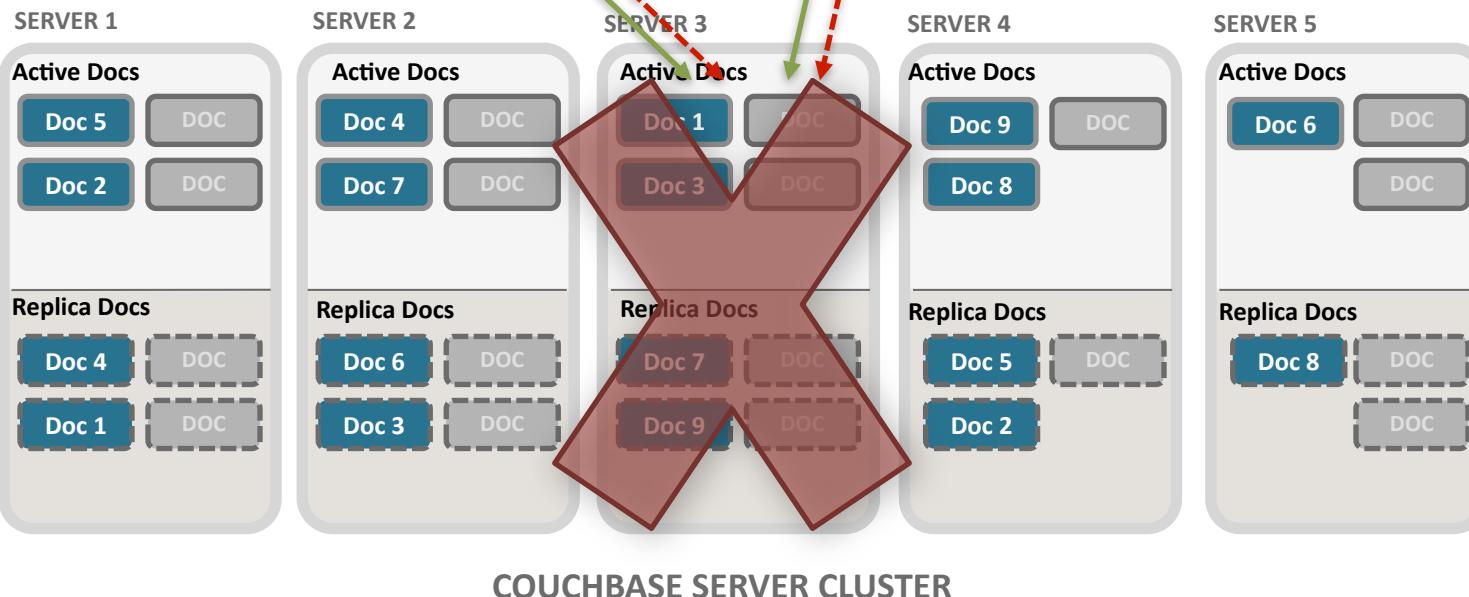
- App servers happily accessing docs on Server 3
- Server fails
- App server requests to server 3 fail
- Cluster detects server has failed
 - Promotes replicas of docs to active
 - Updates cluster map
- App server requests for docs now go to appropriate server
- Typically rebalance would follow



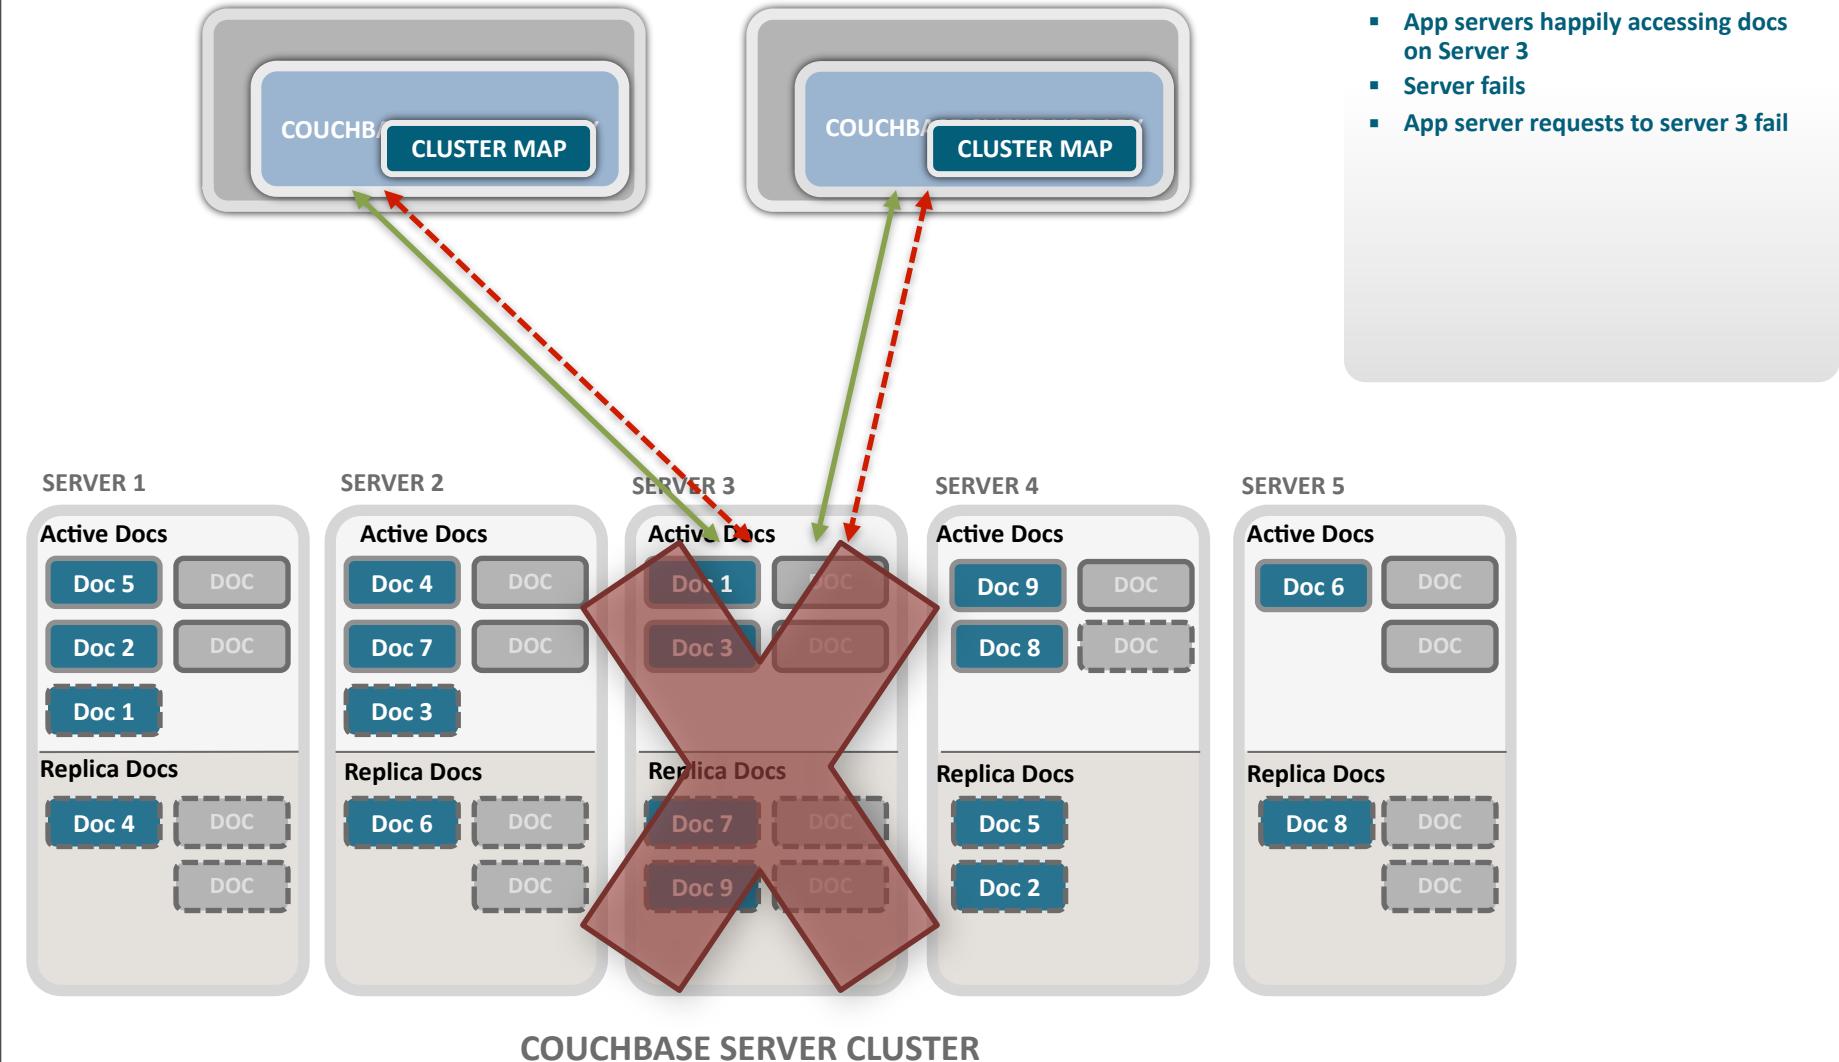
Fail Over Node



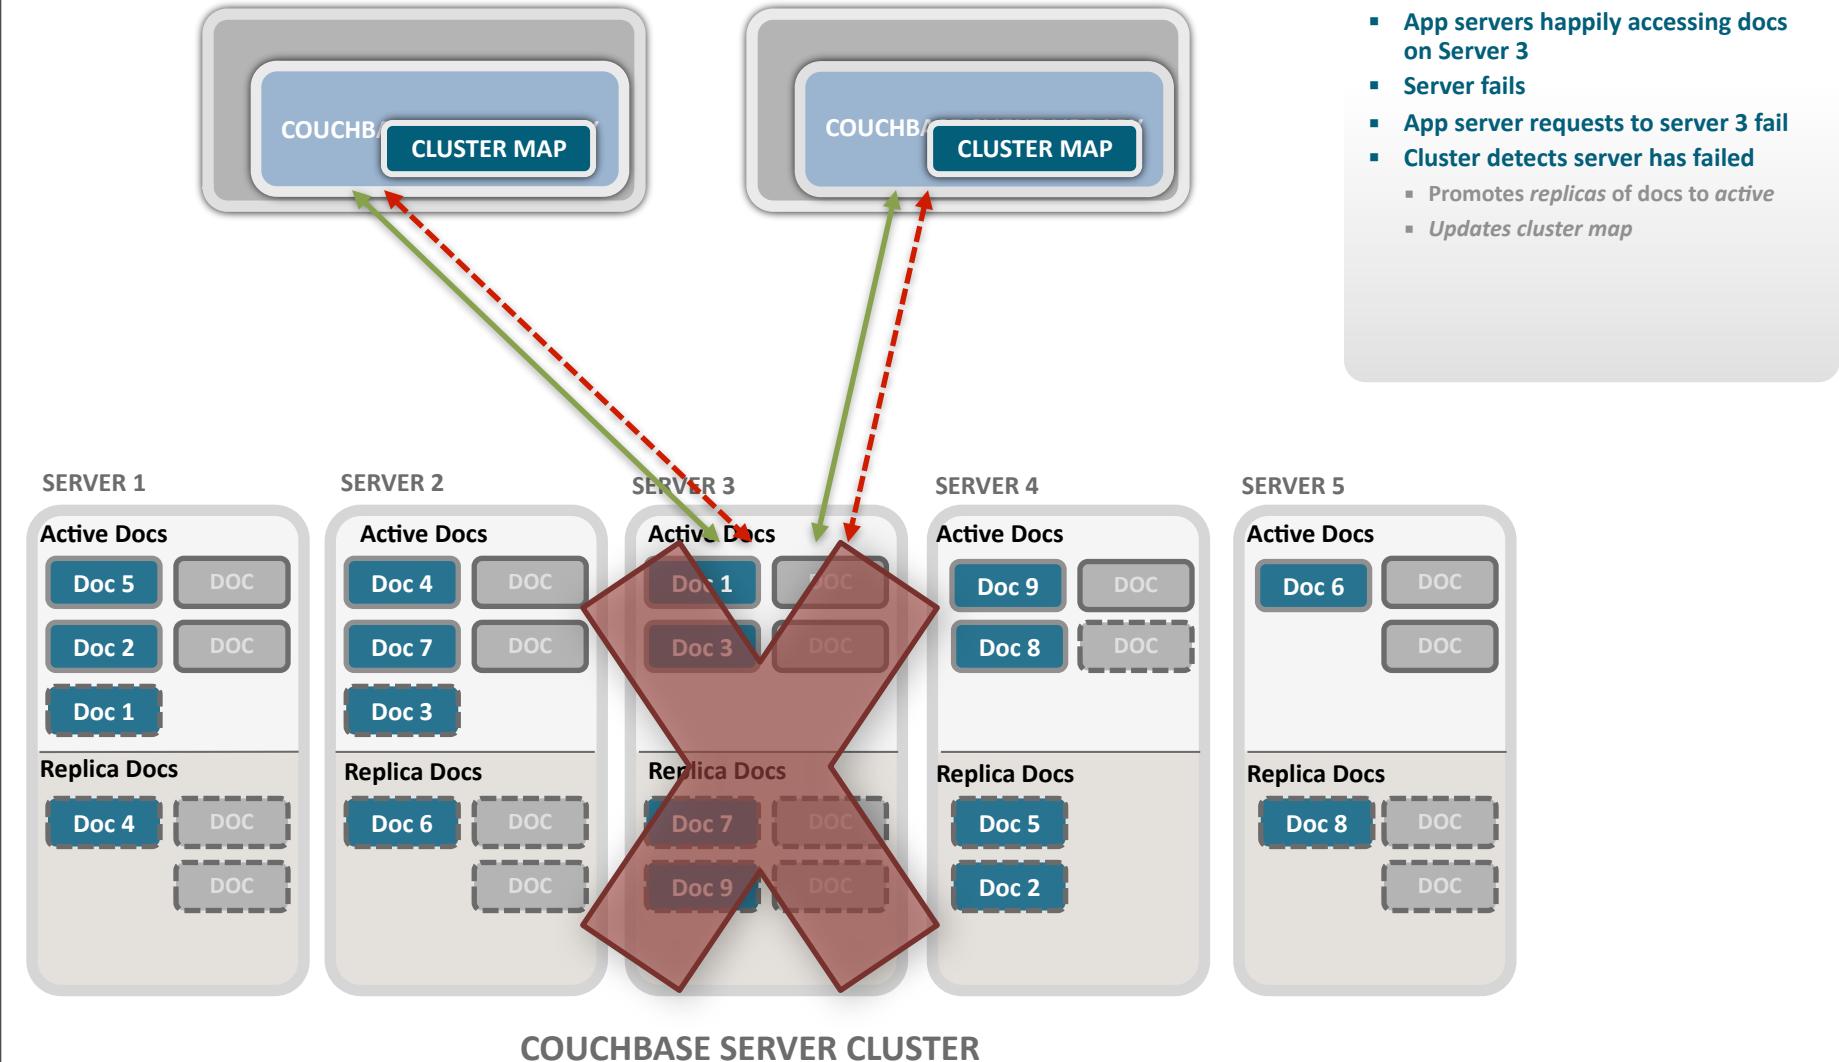
- App servers happily accessing docs on Server 3
- Server fails
- App server requests to server 3 fail
- Cluster detects server has failed
 - Promotes replicas of docs to active
 - Updates cluster map
- App server requests for docs now go to appropriate server
- Typically rebalance would follow



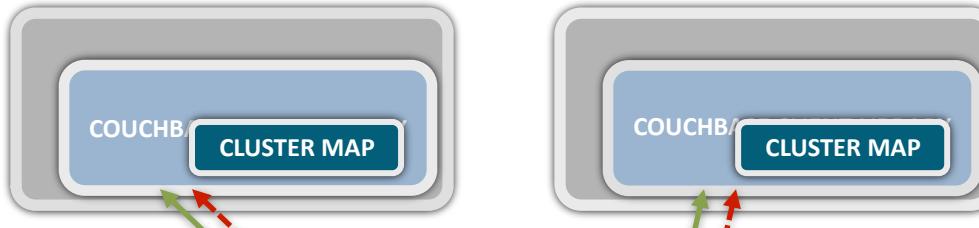
Fail Over Node



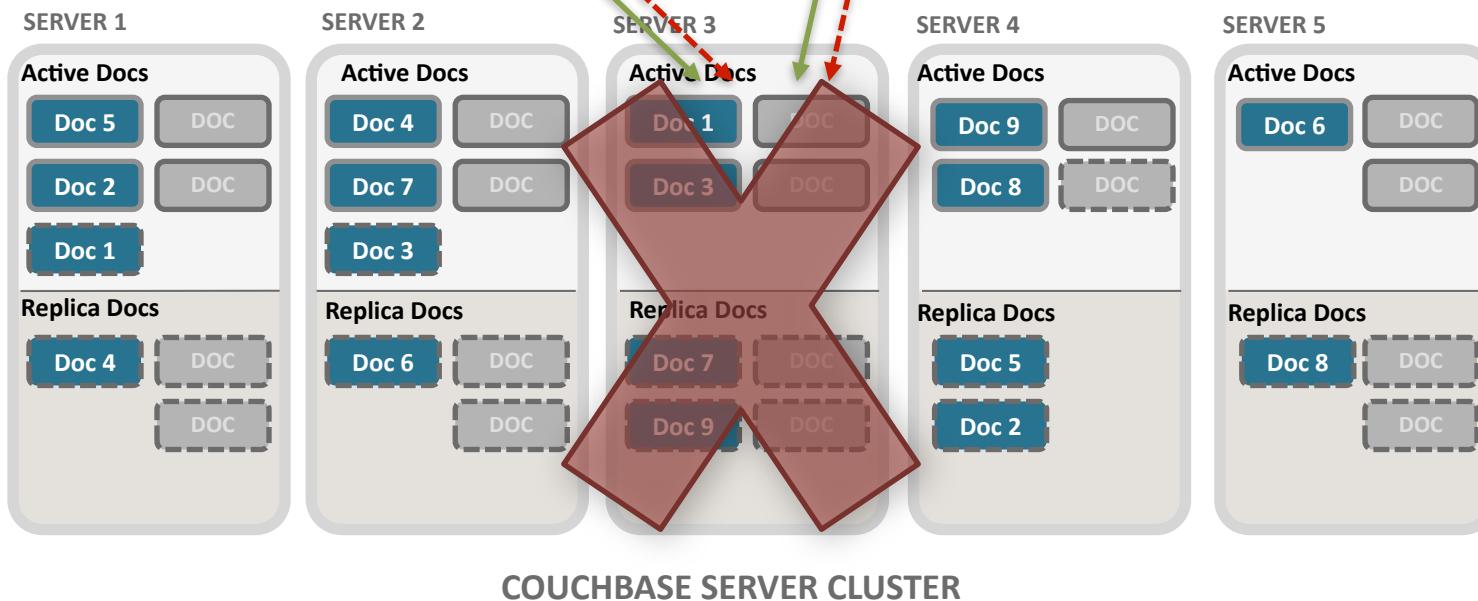
Fail Over Node



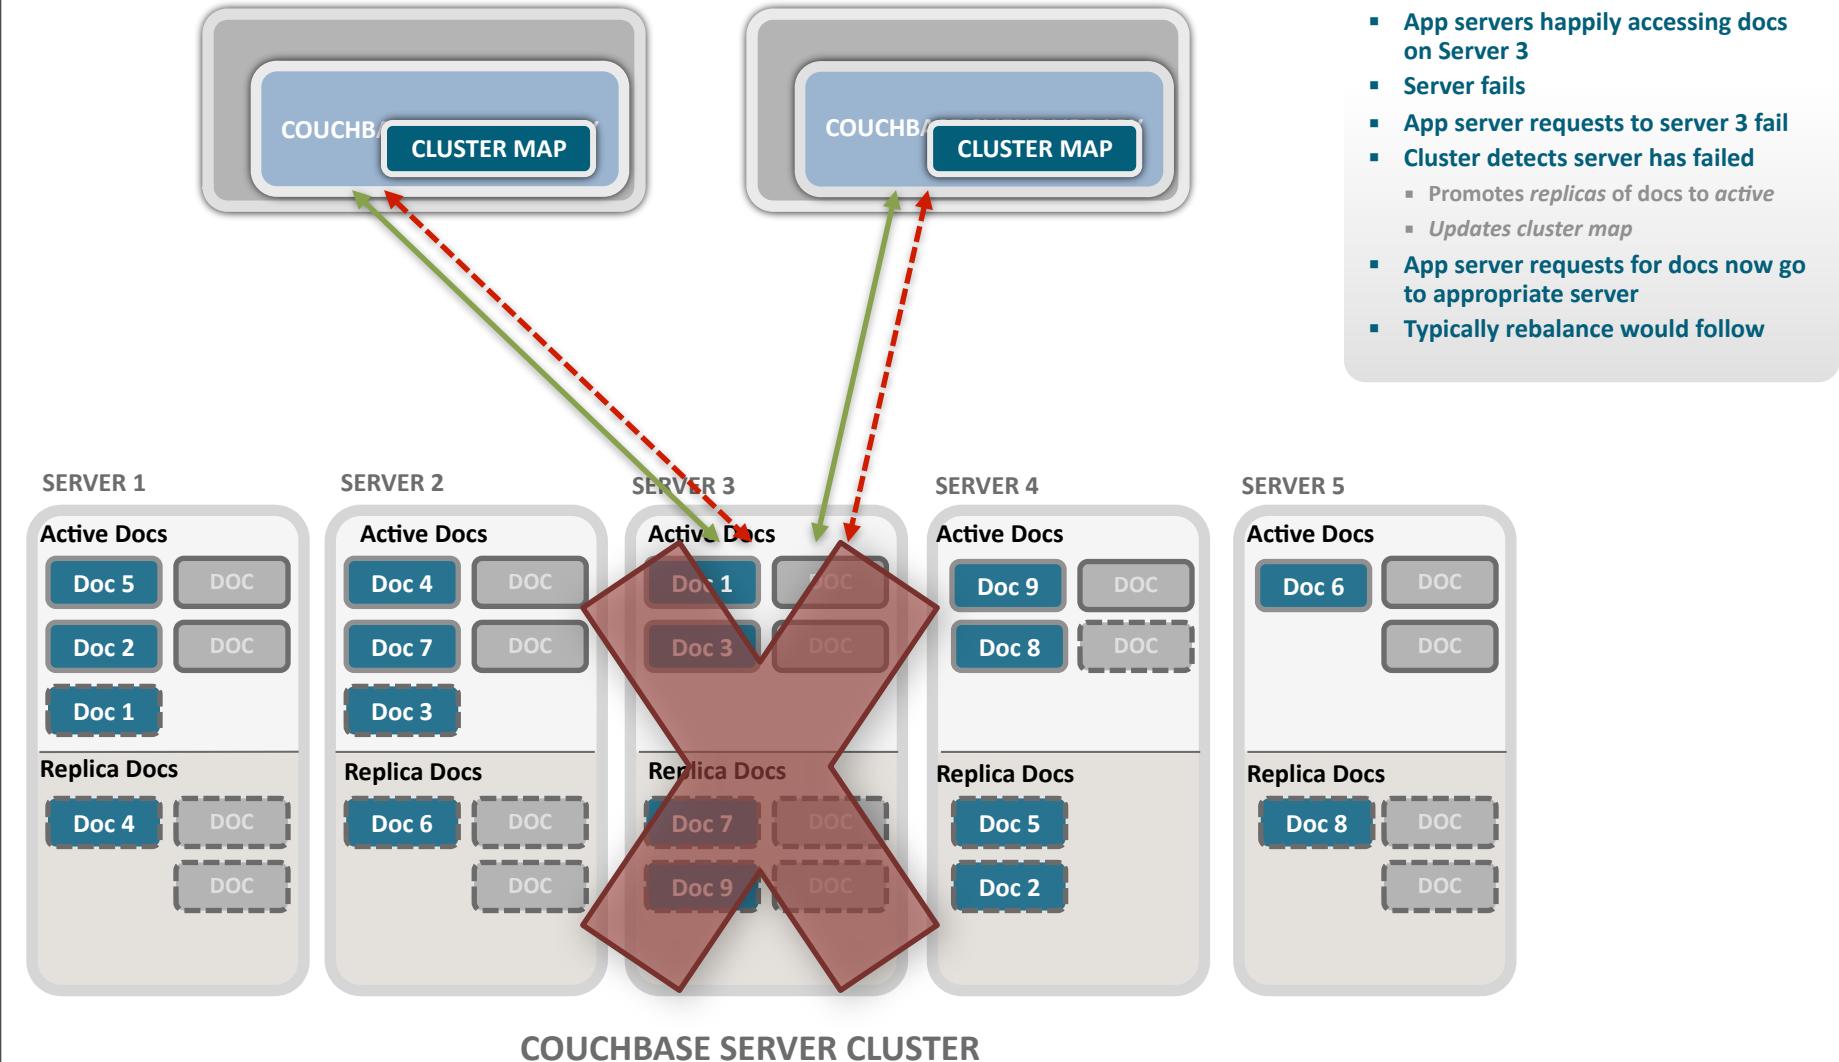
Fail Over Node



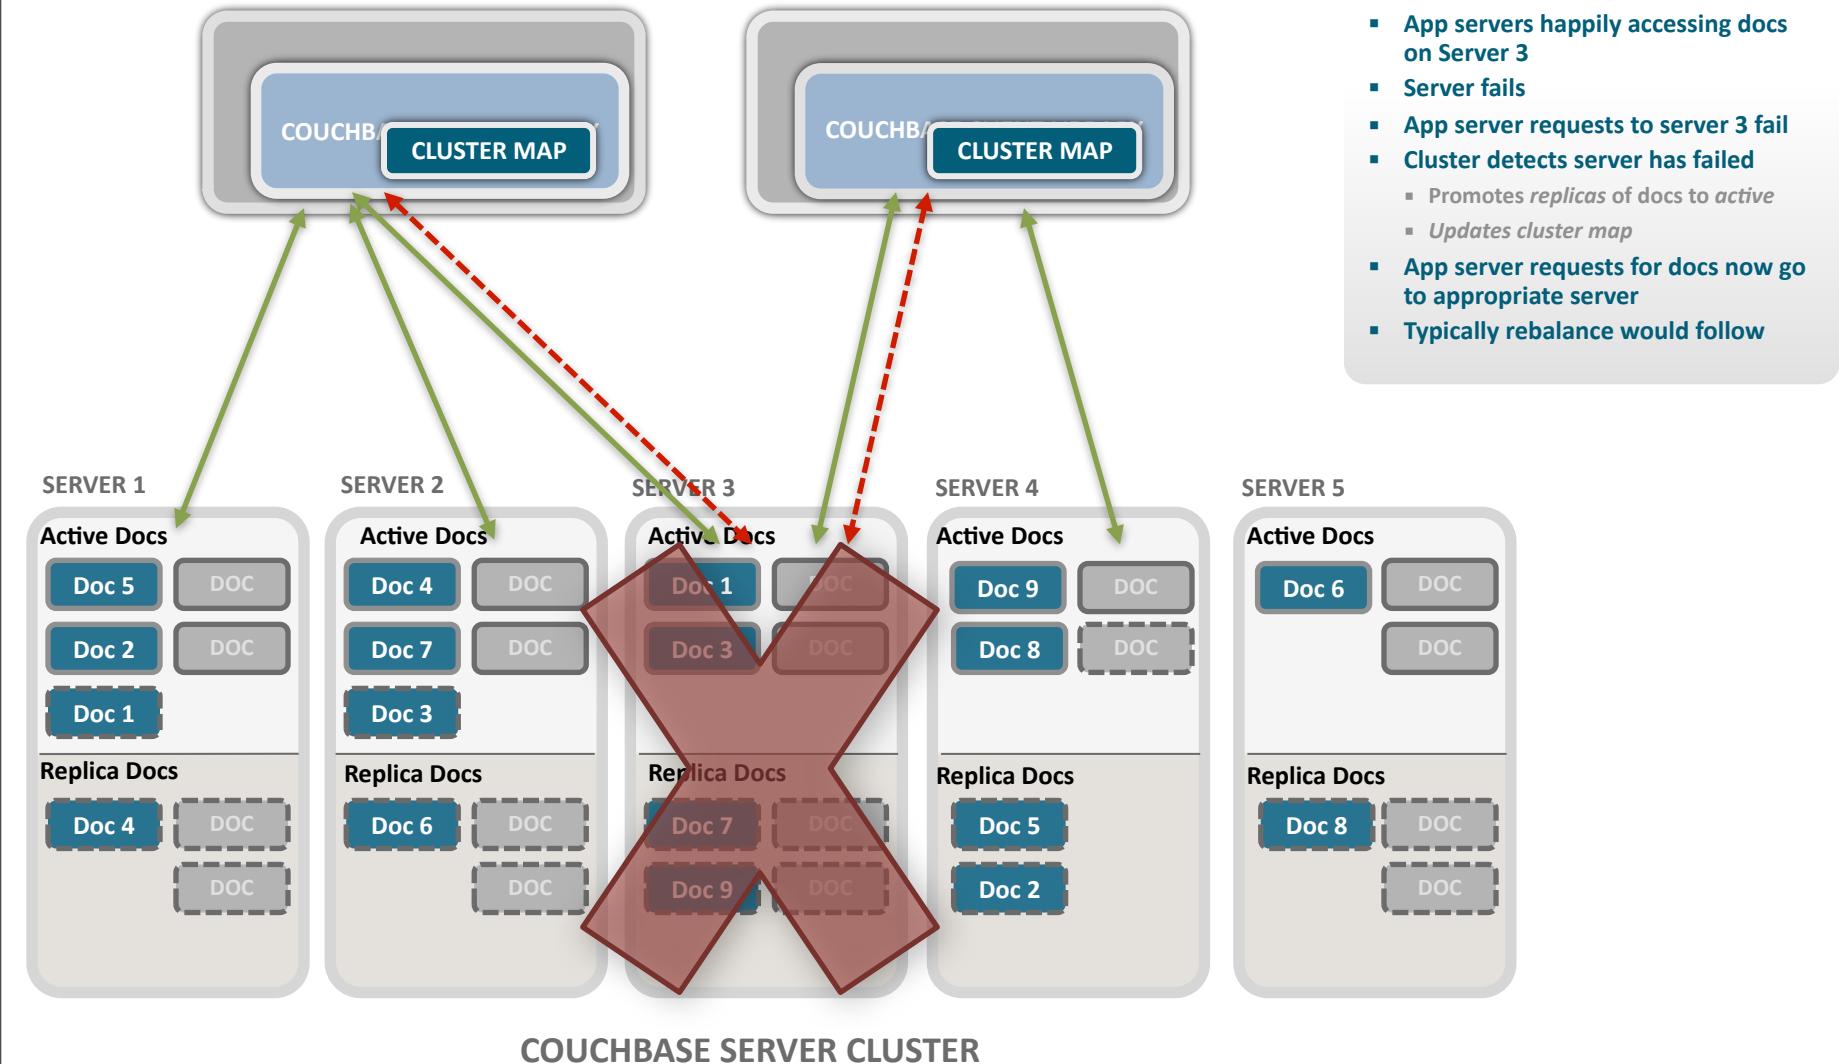
- App servers happily accessing docs on Server 3
- Server fails
- App server requests to server 3 fail
- Cluster detects server has failed
 - Promotes replicas of docs to active
 - Updates cluster map
- App server requests for docs now go to appropriate server



Fail Over Node



Fail Over Node



- App servers happily accessing docs on Server 3
- Server fails
- App server requests to server 3 fail
- Cluster detects server has failed
 - Promotes replicas of docs to active
 - Updates cluster map
- App server requests for docs now go to appropriate server
- Typically rebalance would follow

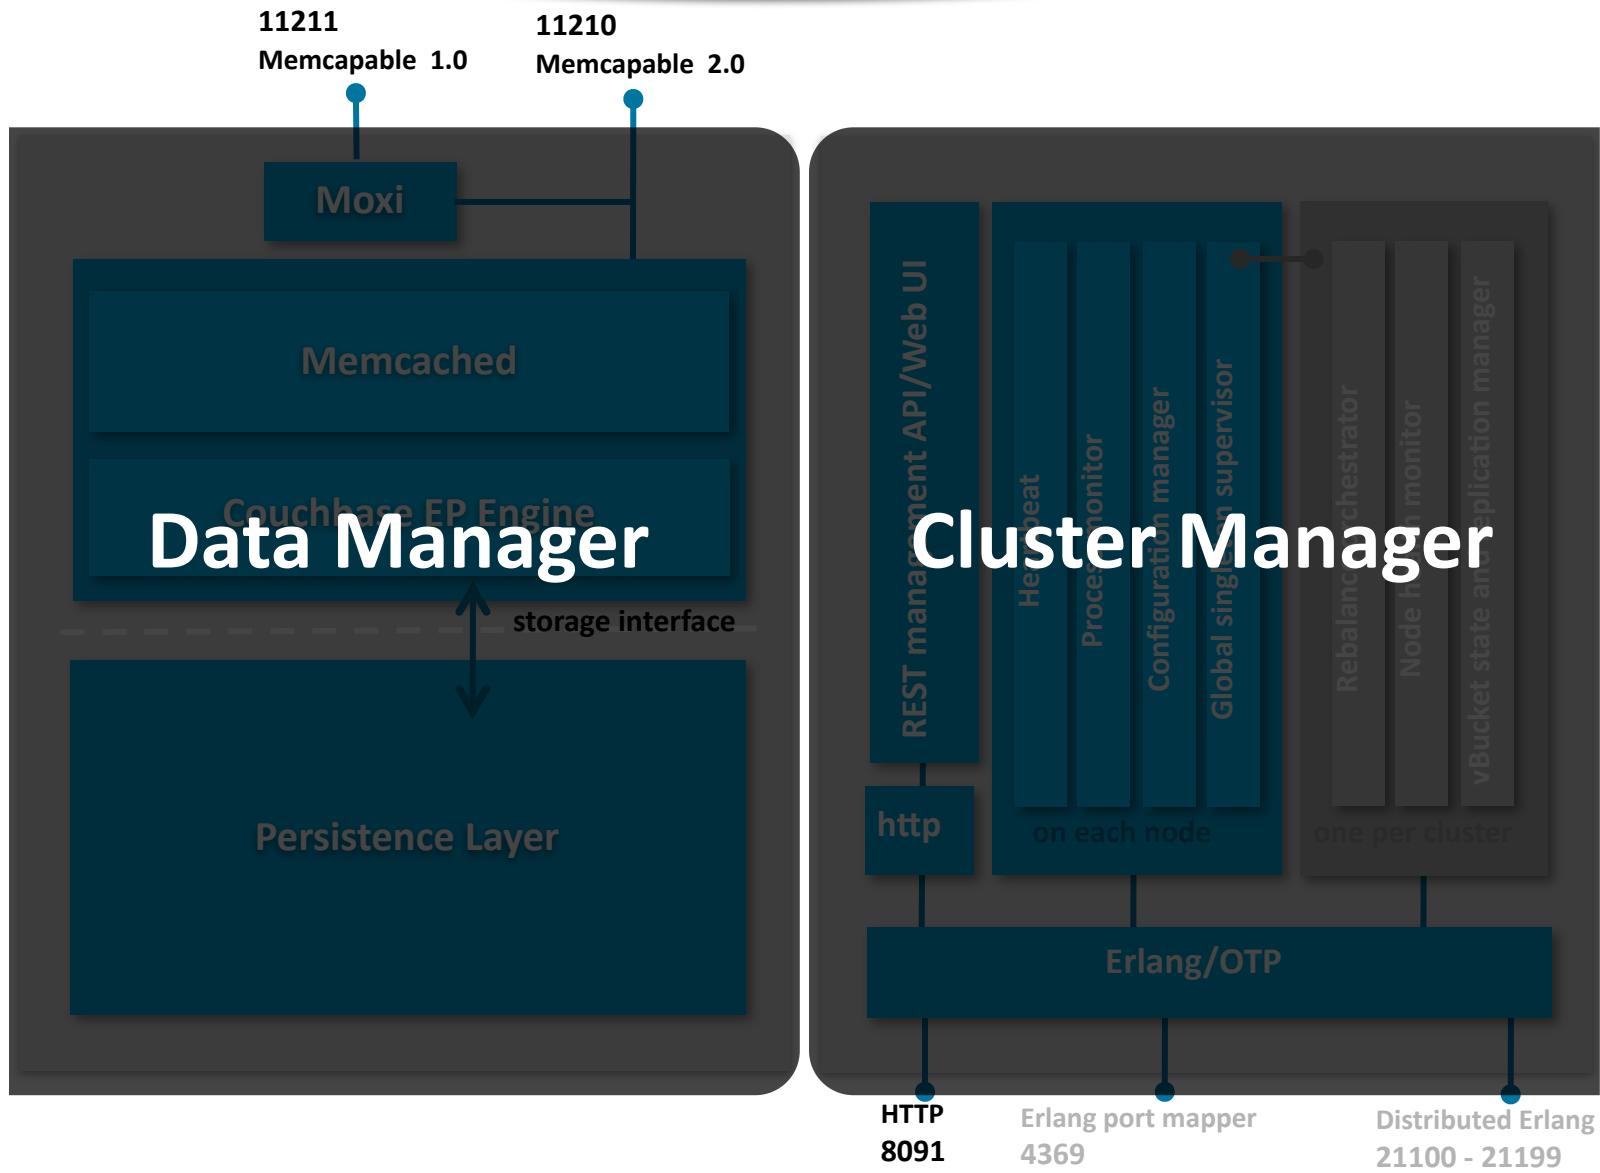
Couchbase Server Features



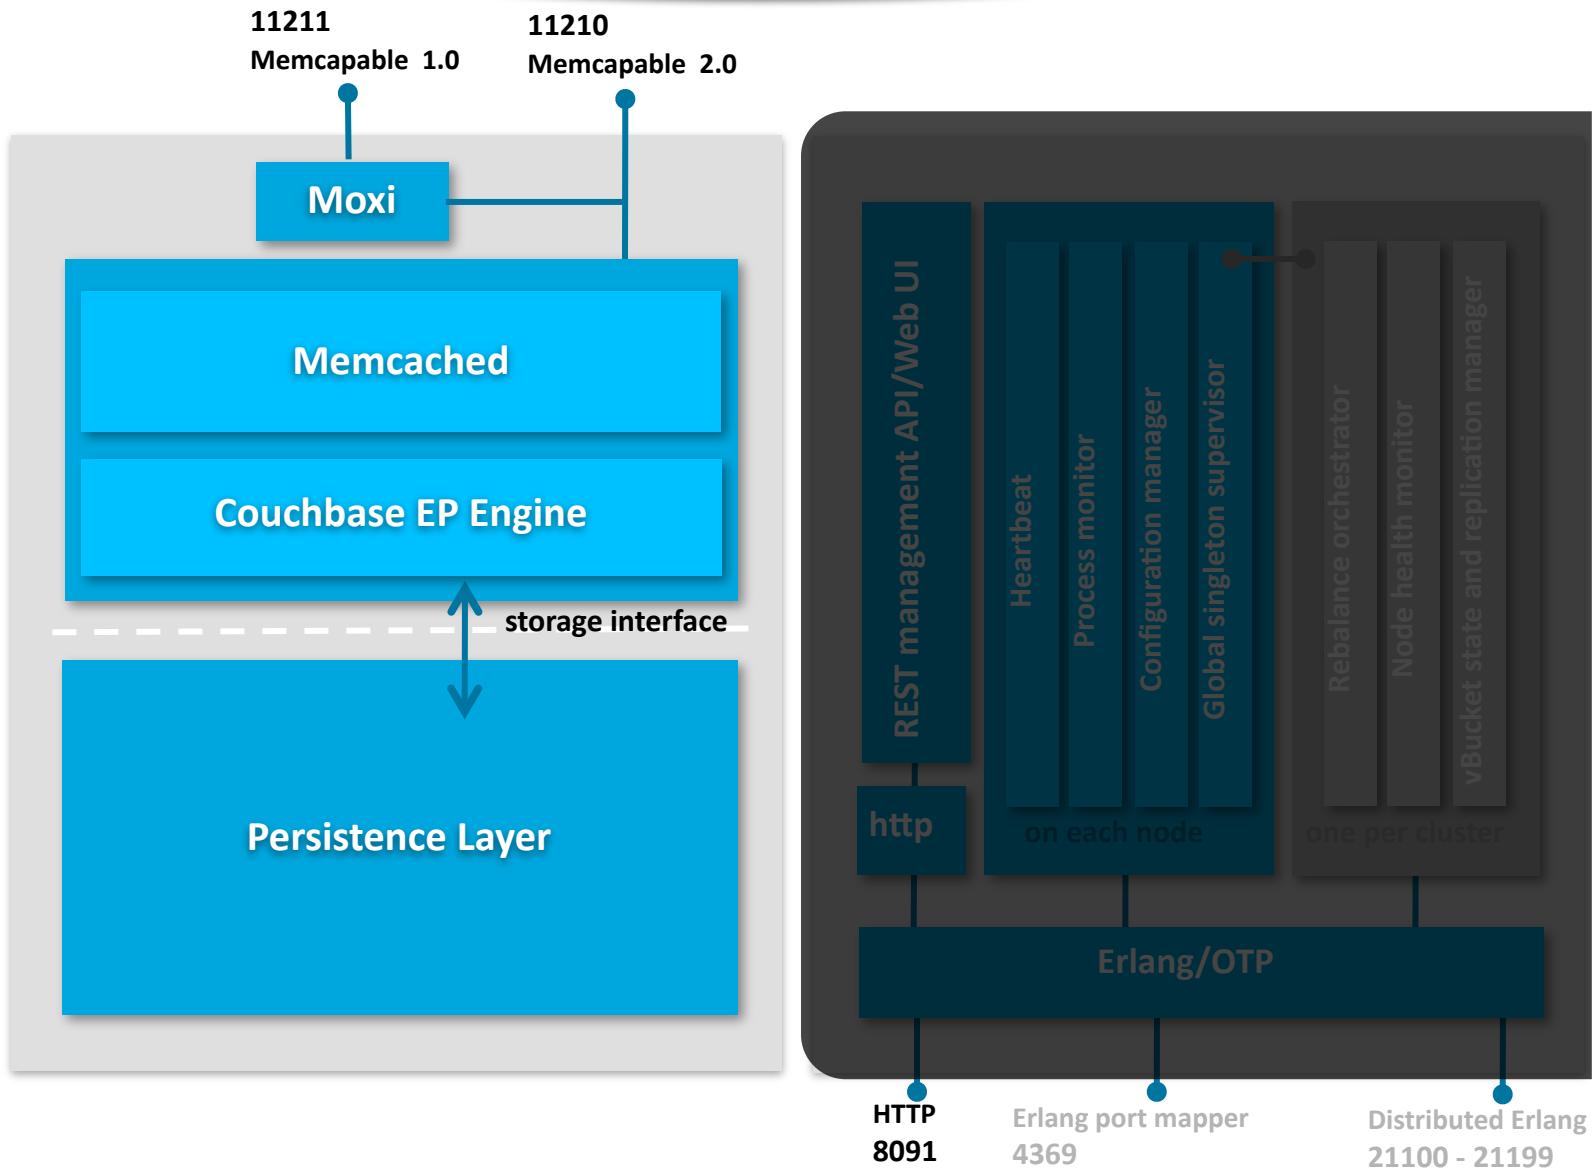
- **Memcached compatible (built-in caching)**
- **Monitoring and administration APIs and GUI**
- **Reliable storage architecture**



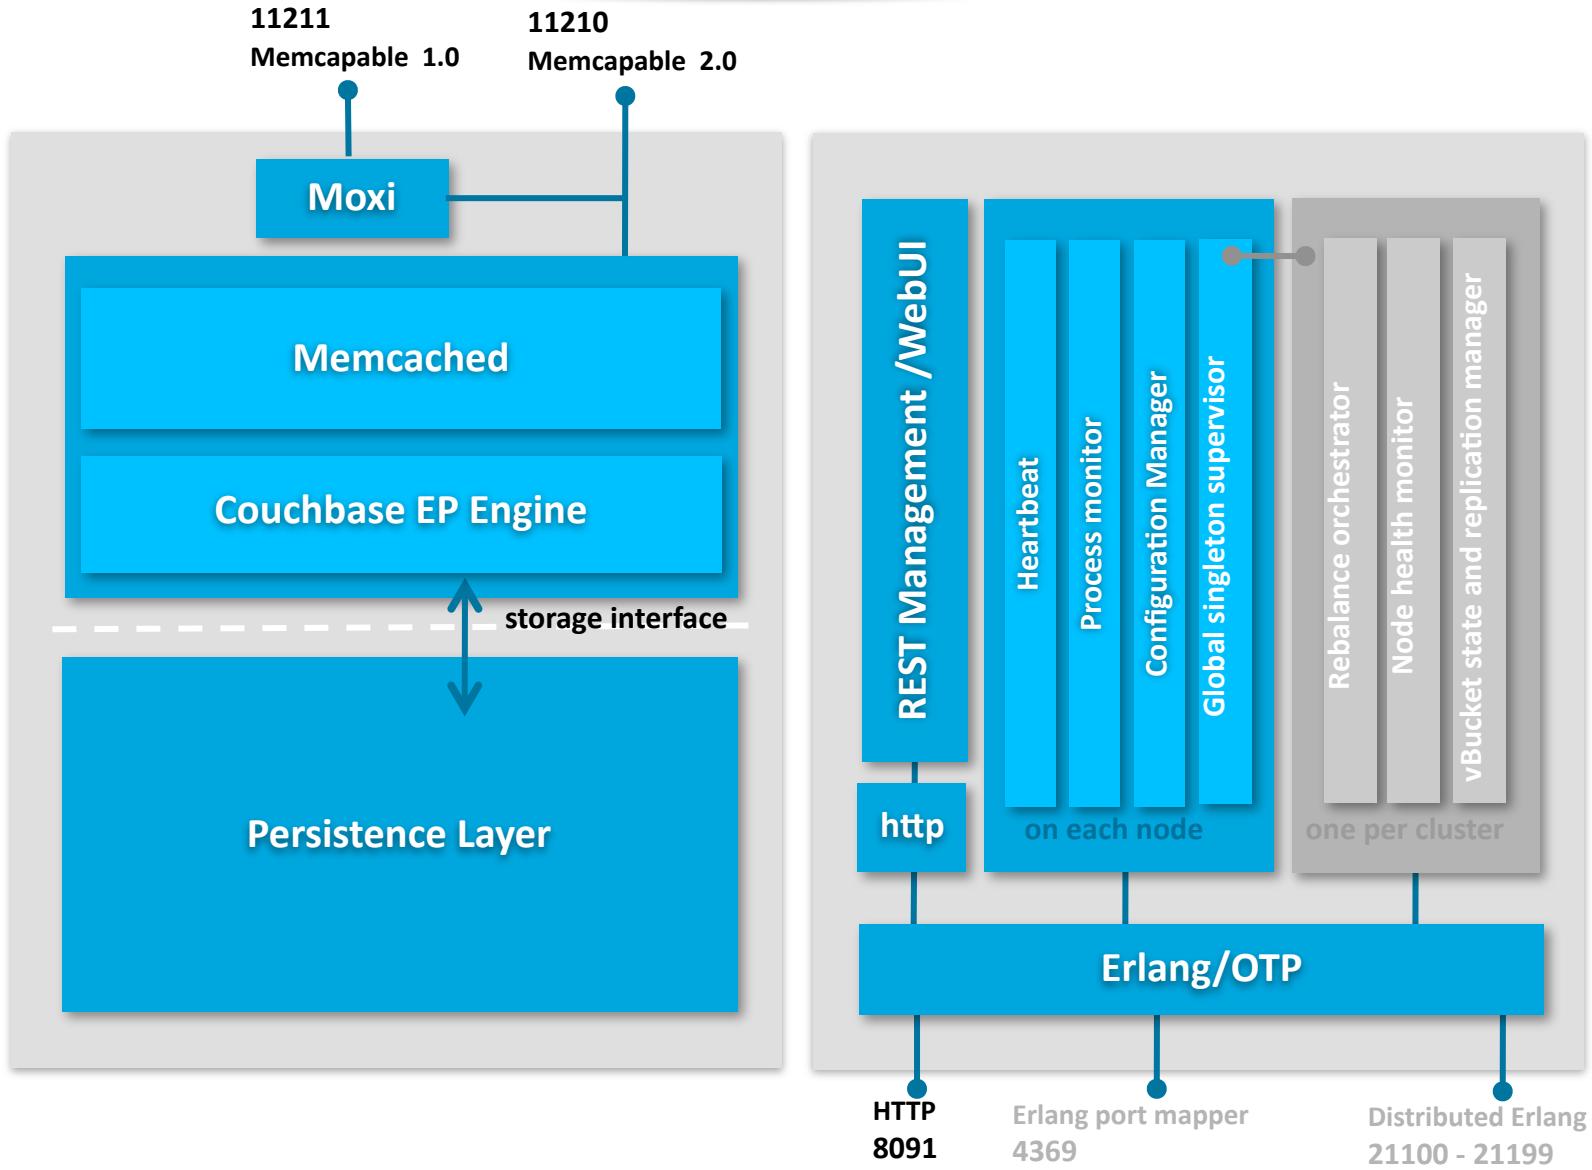
Couchbase Server 1.8 Architecture



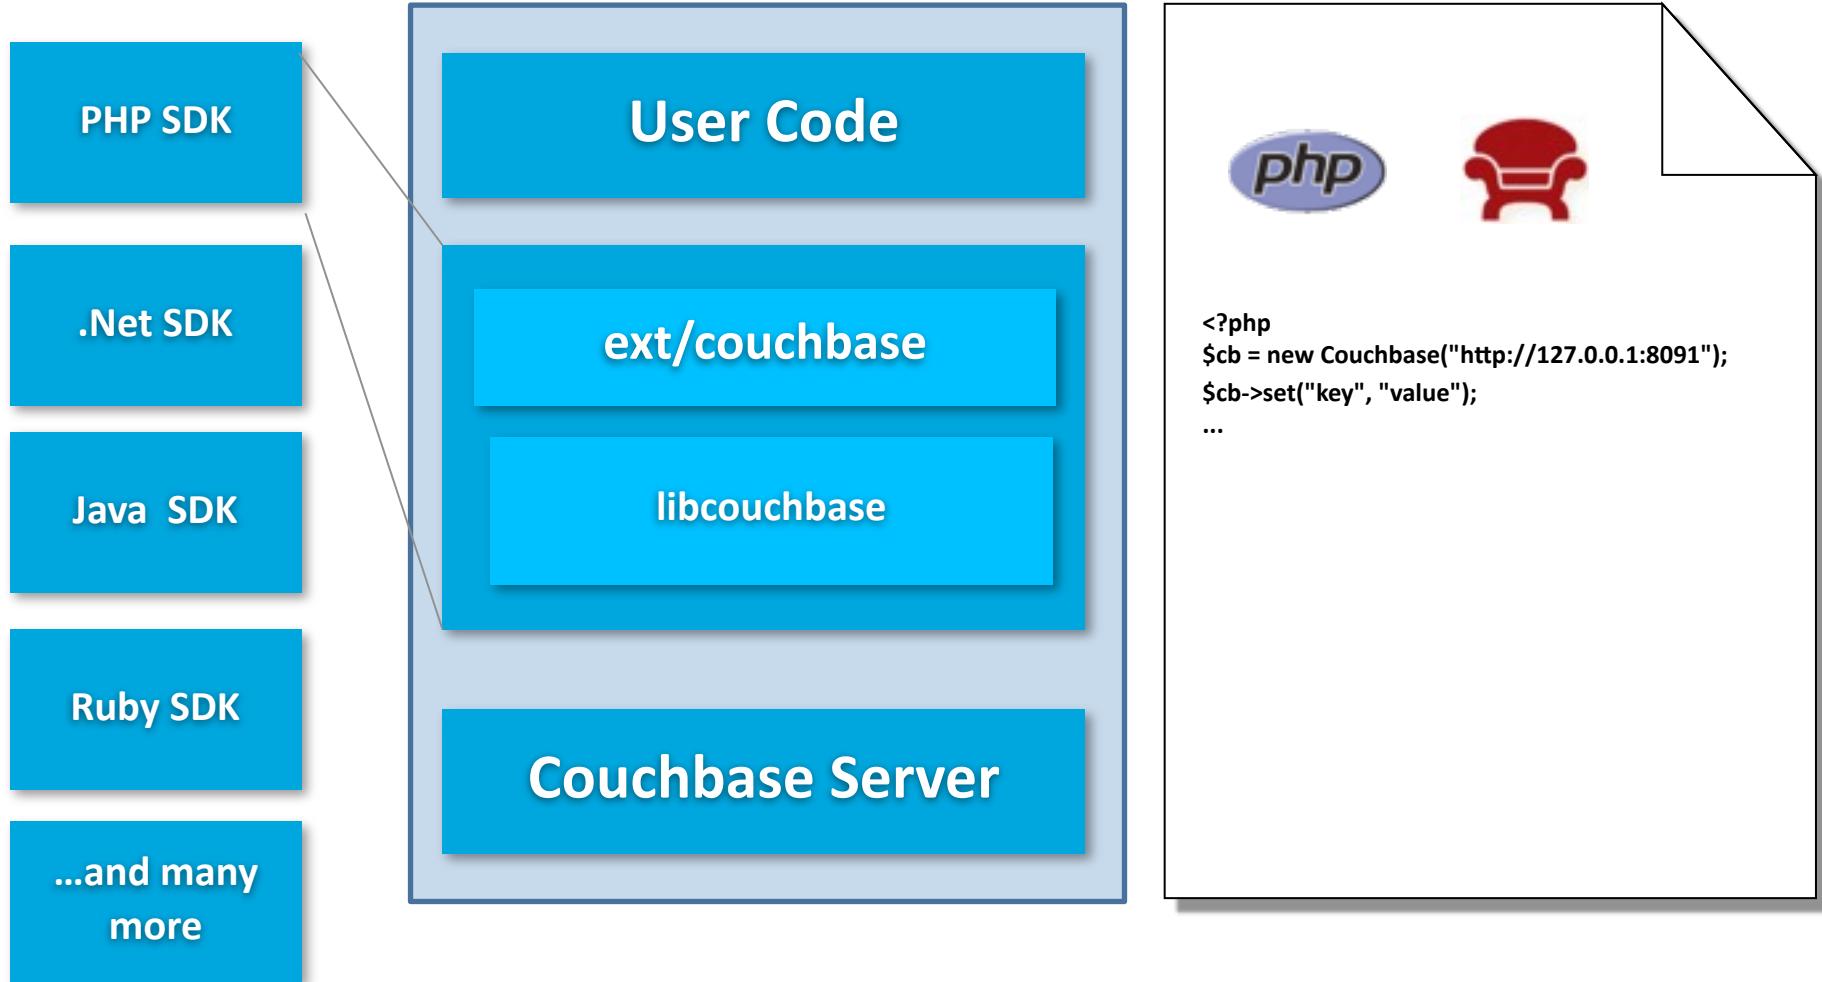
Couchbase Server 1.8 Architecture



Couchbase Server 1.8 Architecture



Couchbase SDKs



<http://www.couchbase.com/develop>

Couchbase Server 2.0

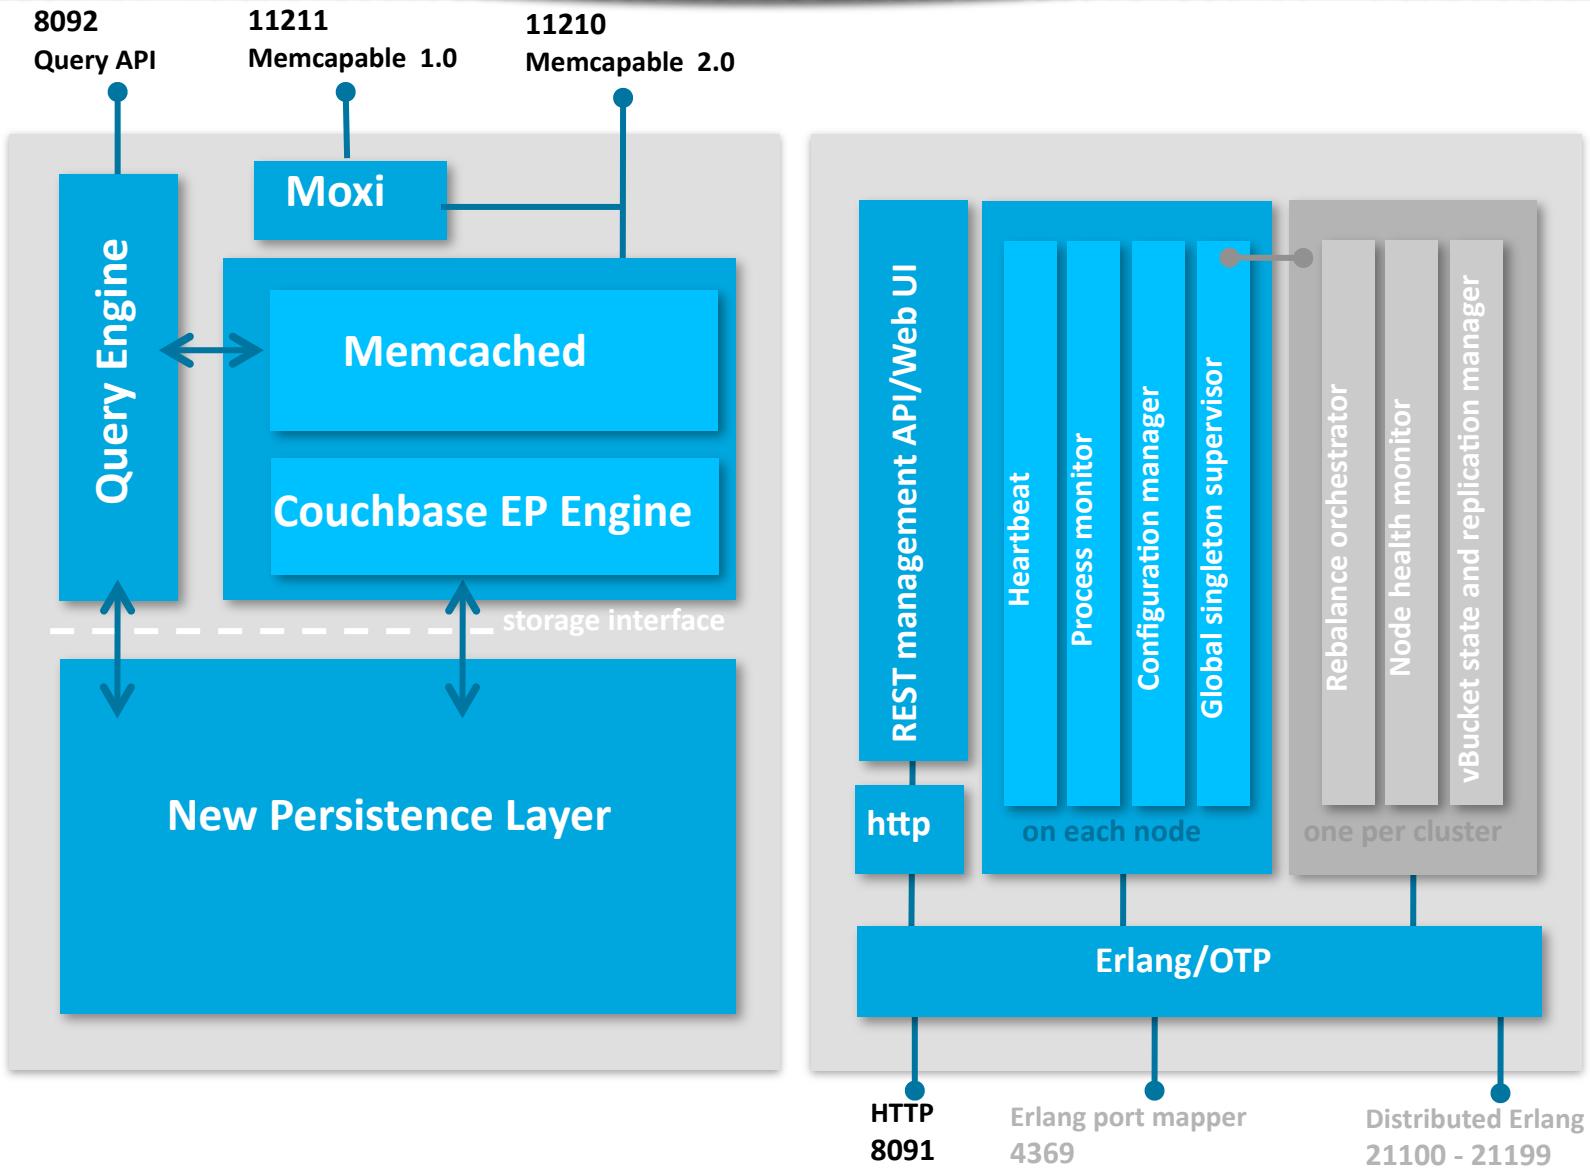
- Next major release of Couchbase Server
- Currently in Developer Preview

What's new:

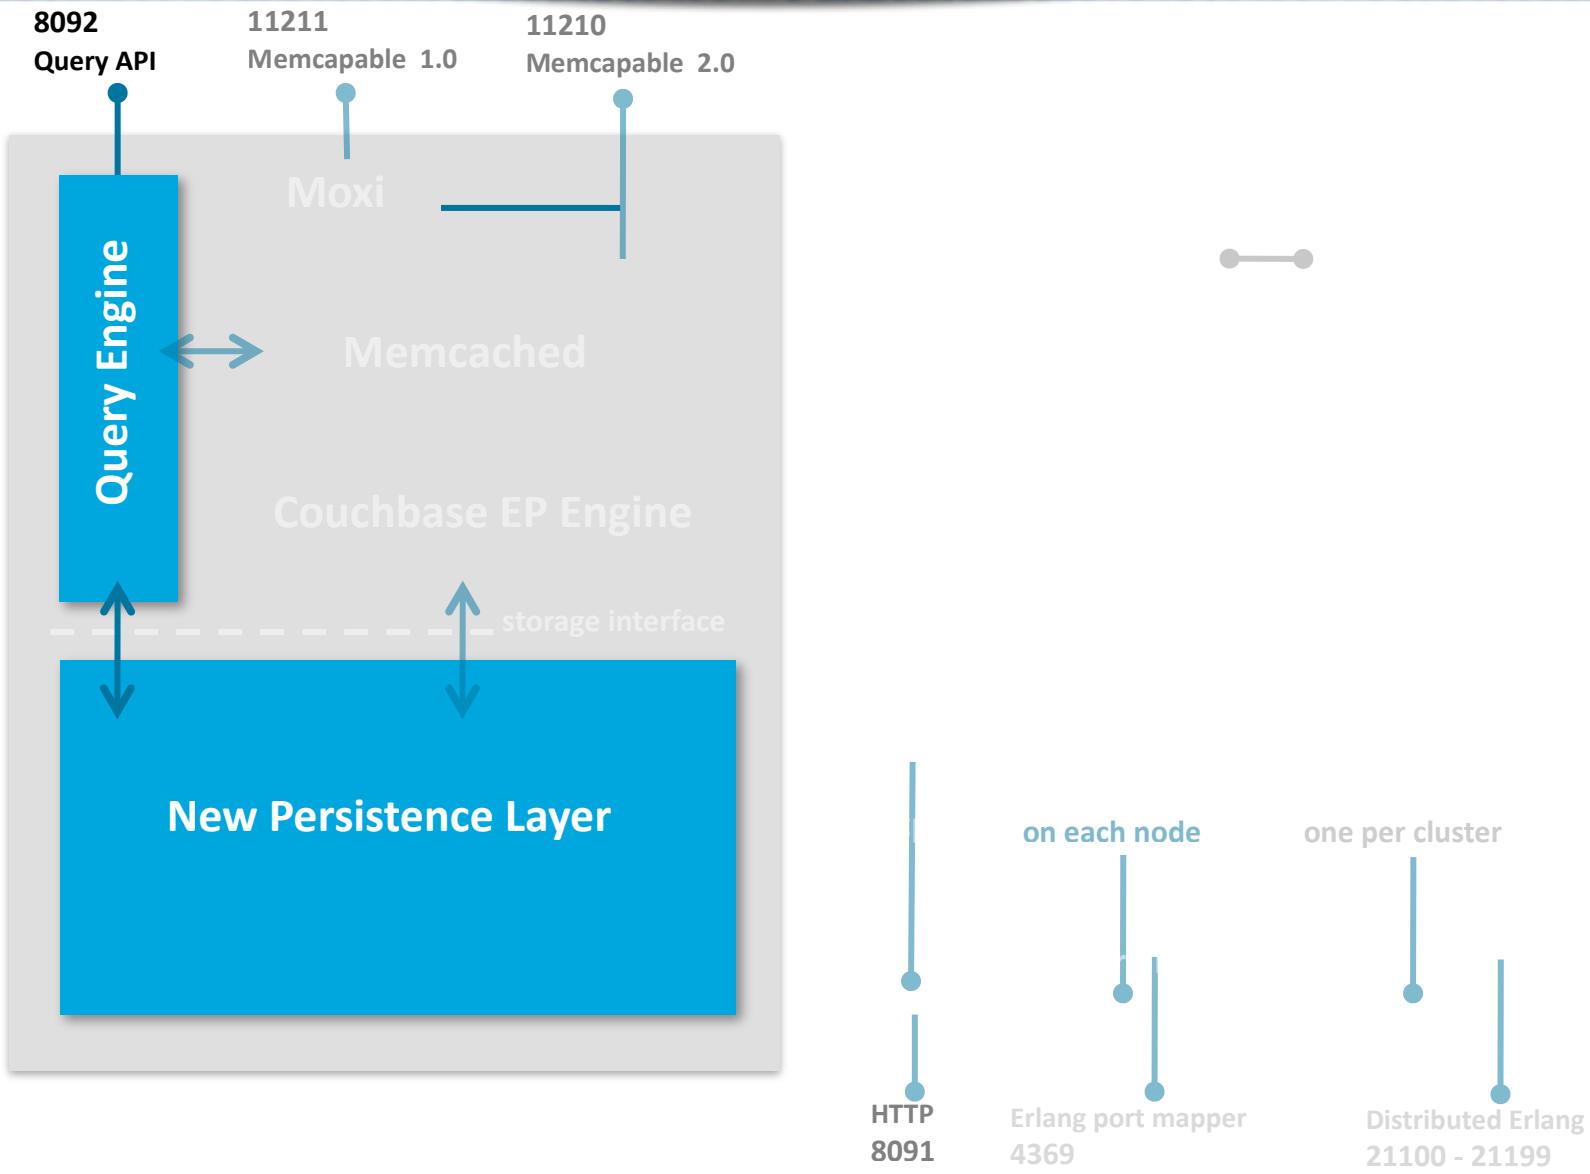
- Indexing and Querying
- Incremental Map Reduce
- Cross Data Center Replication
- Fully backwards compatible with existing Couchbase Server



Couchbase Server 2.0 Architecture



Couchbase Server 2.0 Architecture



Indexing and Querying

APP SERVER 1



APP SERVER 2



- Indexing work is distributed amongst nodes
 - Large data set possible
 - Parallelize the effort

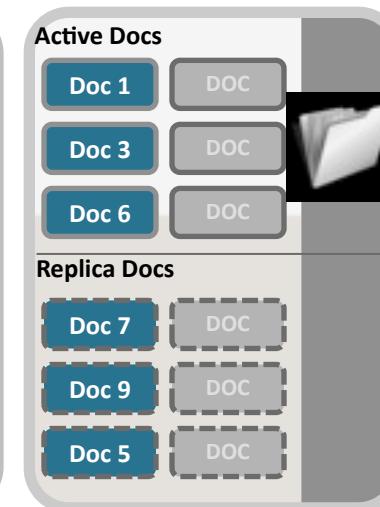
SERVER 1



SERVER 2



SERVER 3



Indexing and Querying

APP SERVER 1



APP SERVER 2

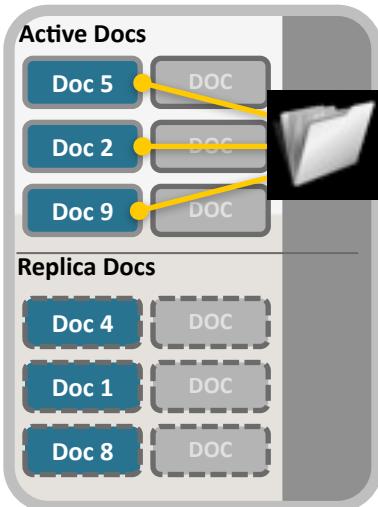


- Indexing work is distributed amongst nodes

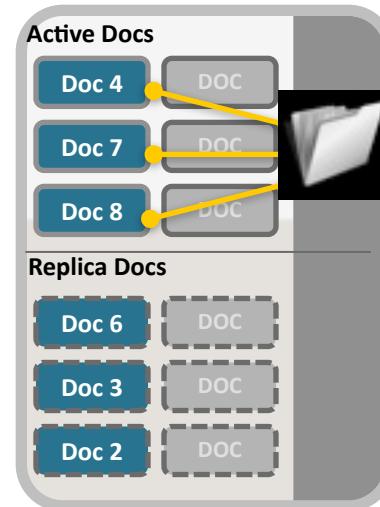
- Large data set possible
- Parallelize the effort

- Each node has index for data stored on it

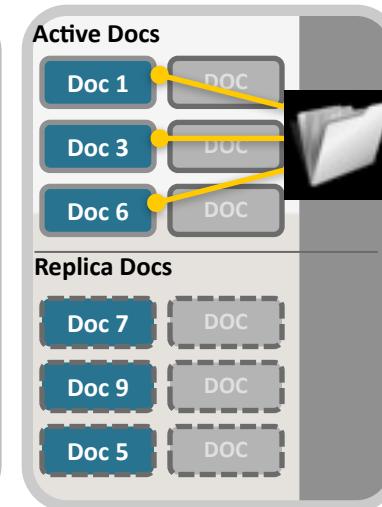
SERVER 1



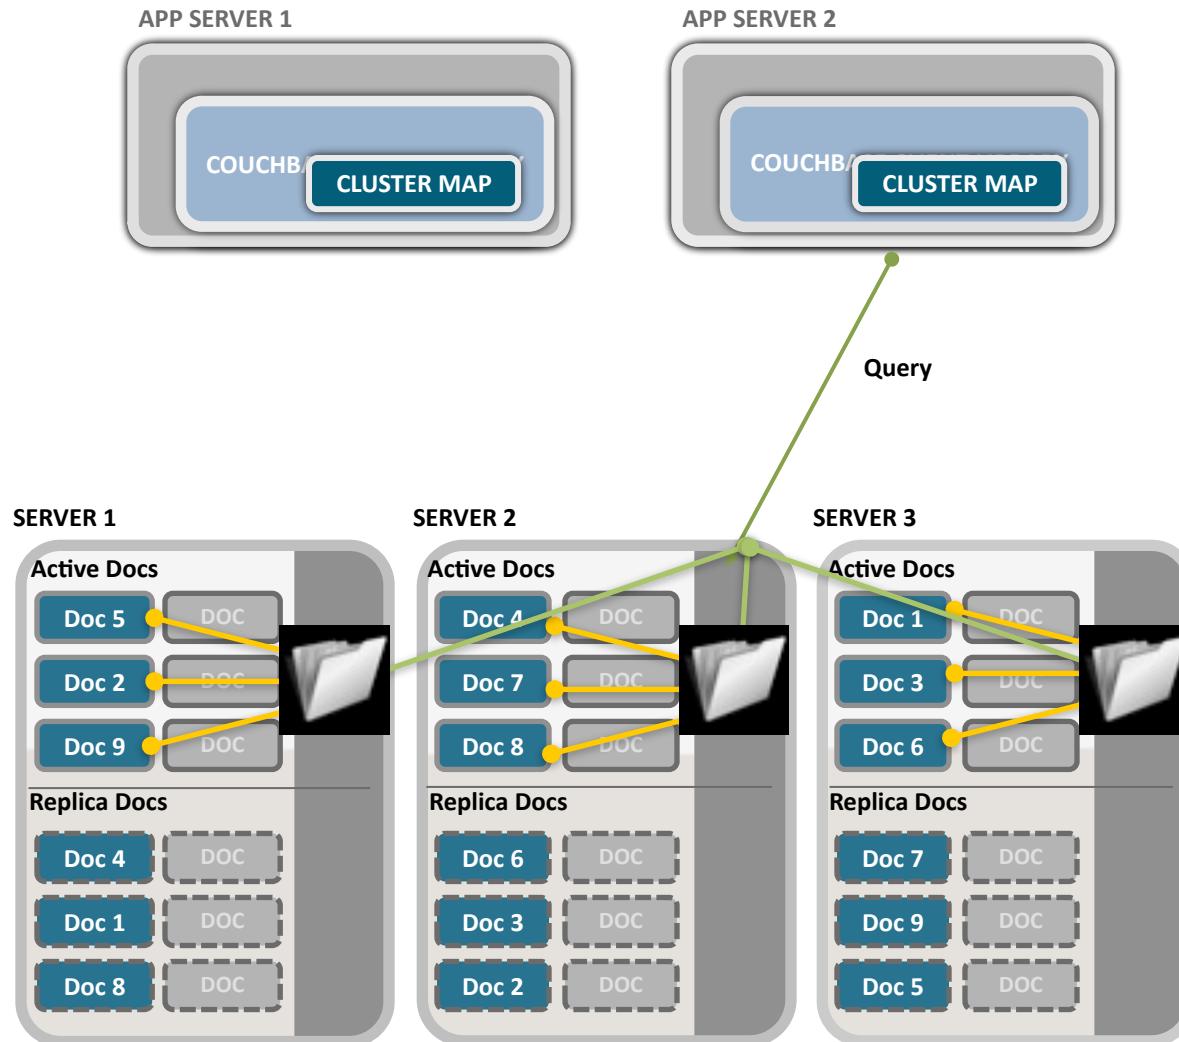
SERVER 2



SERVER 3

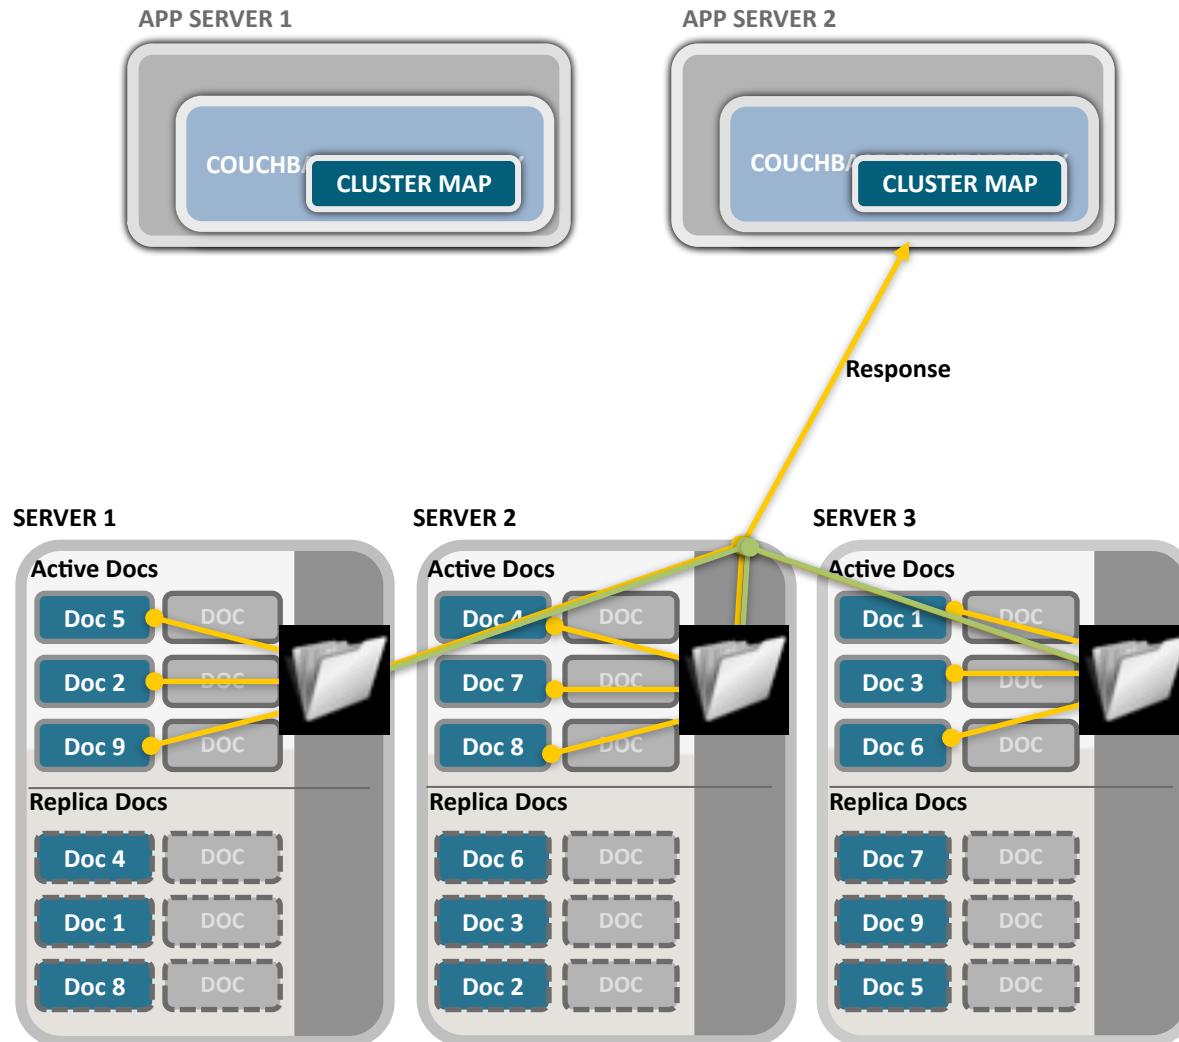


Indexing and Querying



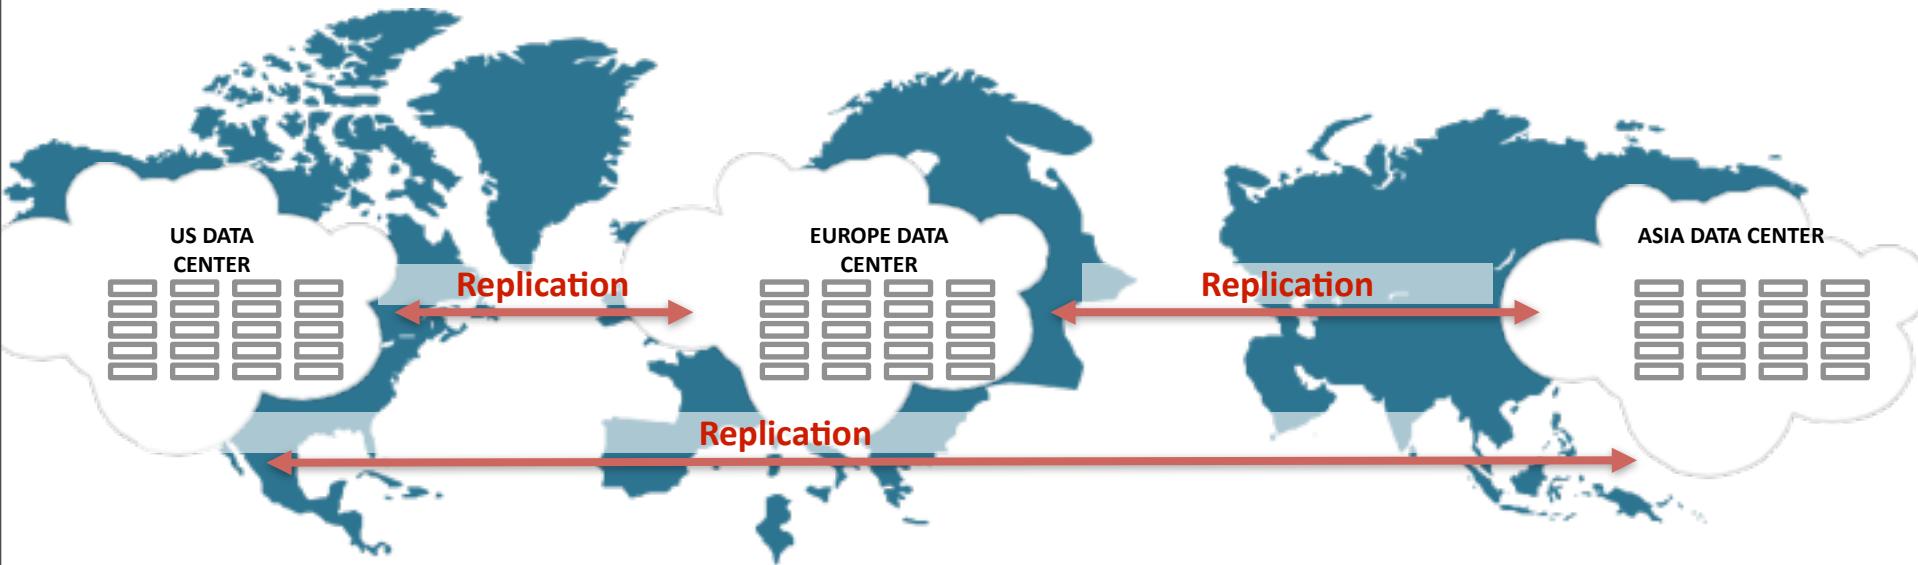
- Indexing work is distributed amongst nodes
 - Large data set possible
 - Parallelize the effort
- Each node has index for data stored on it

Indexing and Querying



- Indexing work is distributed amongst nodes
 - Large data set possible
 - Parallelize the effort
- Each node has index for data stored on it
- Queries combine the results from required nodes

Cross Data Center Replication



- Want data close to user
- Want multiple locations for disaster recovery
- Multi-Master: Can write to same document in all different regions & it will sync (eventually consistent, always available)

EXT/COUCHBASE

Setup

```
<?php  
$url = "http://localhost:8091/";  
$cb = new Couchbase($url);
```

Basics

```
<?php
// setup
$cb->set("a", 1);

$a = $cb->get("a");
echo $a; // prints 1

$cb->increment("a");
echo $cb->get("a"); // prints 2
$cb->delete("a"); // booya
```

Storage Operations

```
<?php  
// set "a" to 1  
$cb->set("a", 1);  
  
// fails if "b" exists  
$cb->add("b", 1);  
  
// fails if "c" doesn't exist  
$cb->replace("c", 1);
```

Expiration

```
<?php
// set "a" to 1 for 10 seconds
$cb->set("a", 1, 10);

// most other ops can set expire

// just update the expiry
$cb->touch("a", 10);
```

Compare And Swap

```
<?php  
$cas = -1;  
$cb->get("a", null, &$cas);  
  
echo $cas; // prints 76324827359  
  
// fails if cas doesn't match  
$cb->cas($cas, "a", 2);
```

Arithmetic

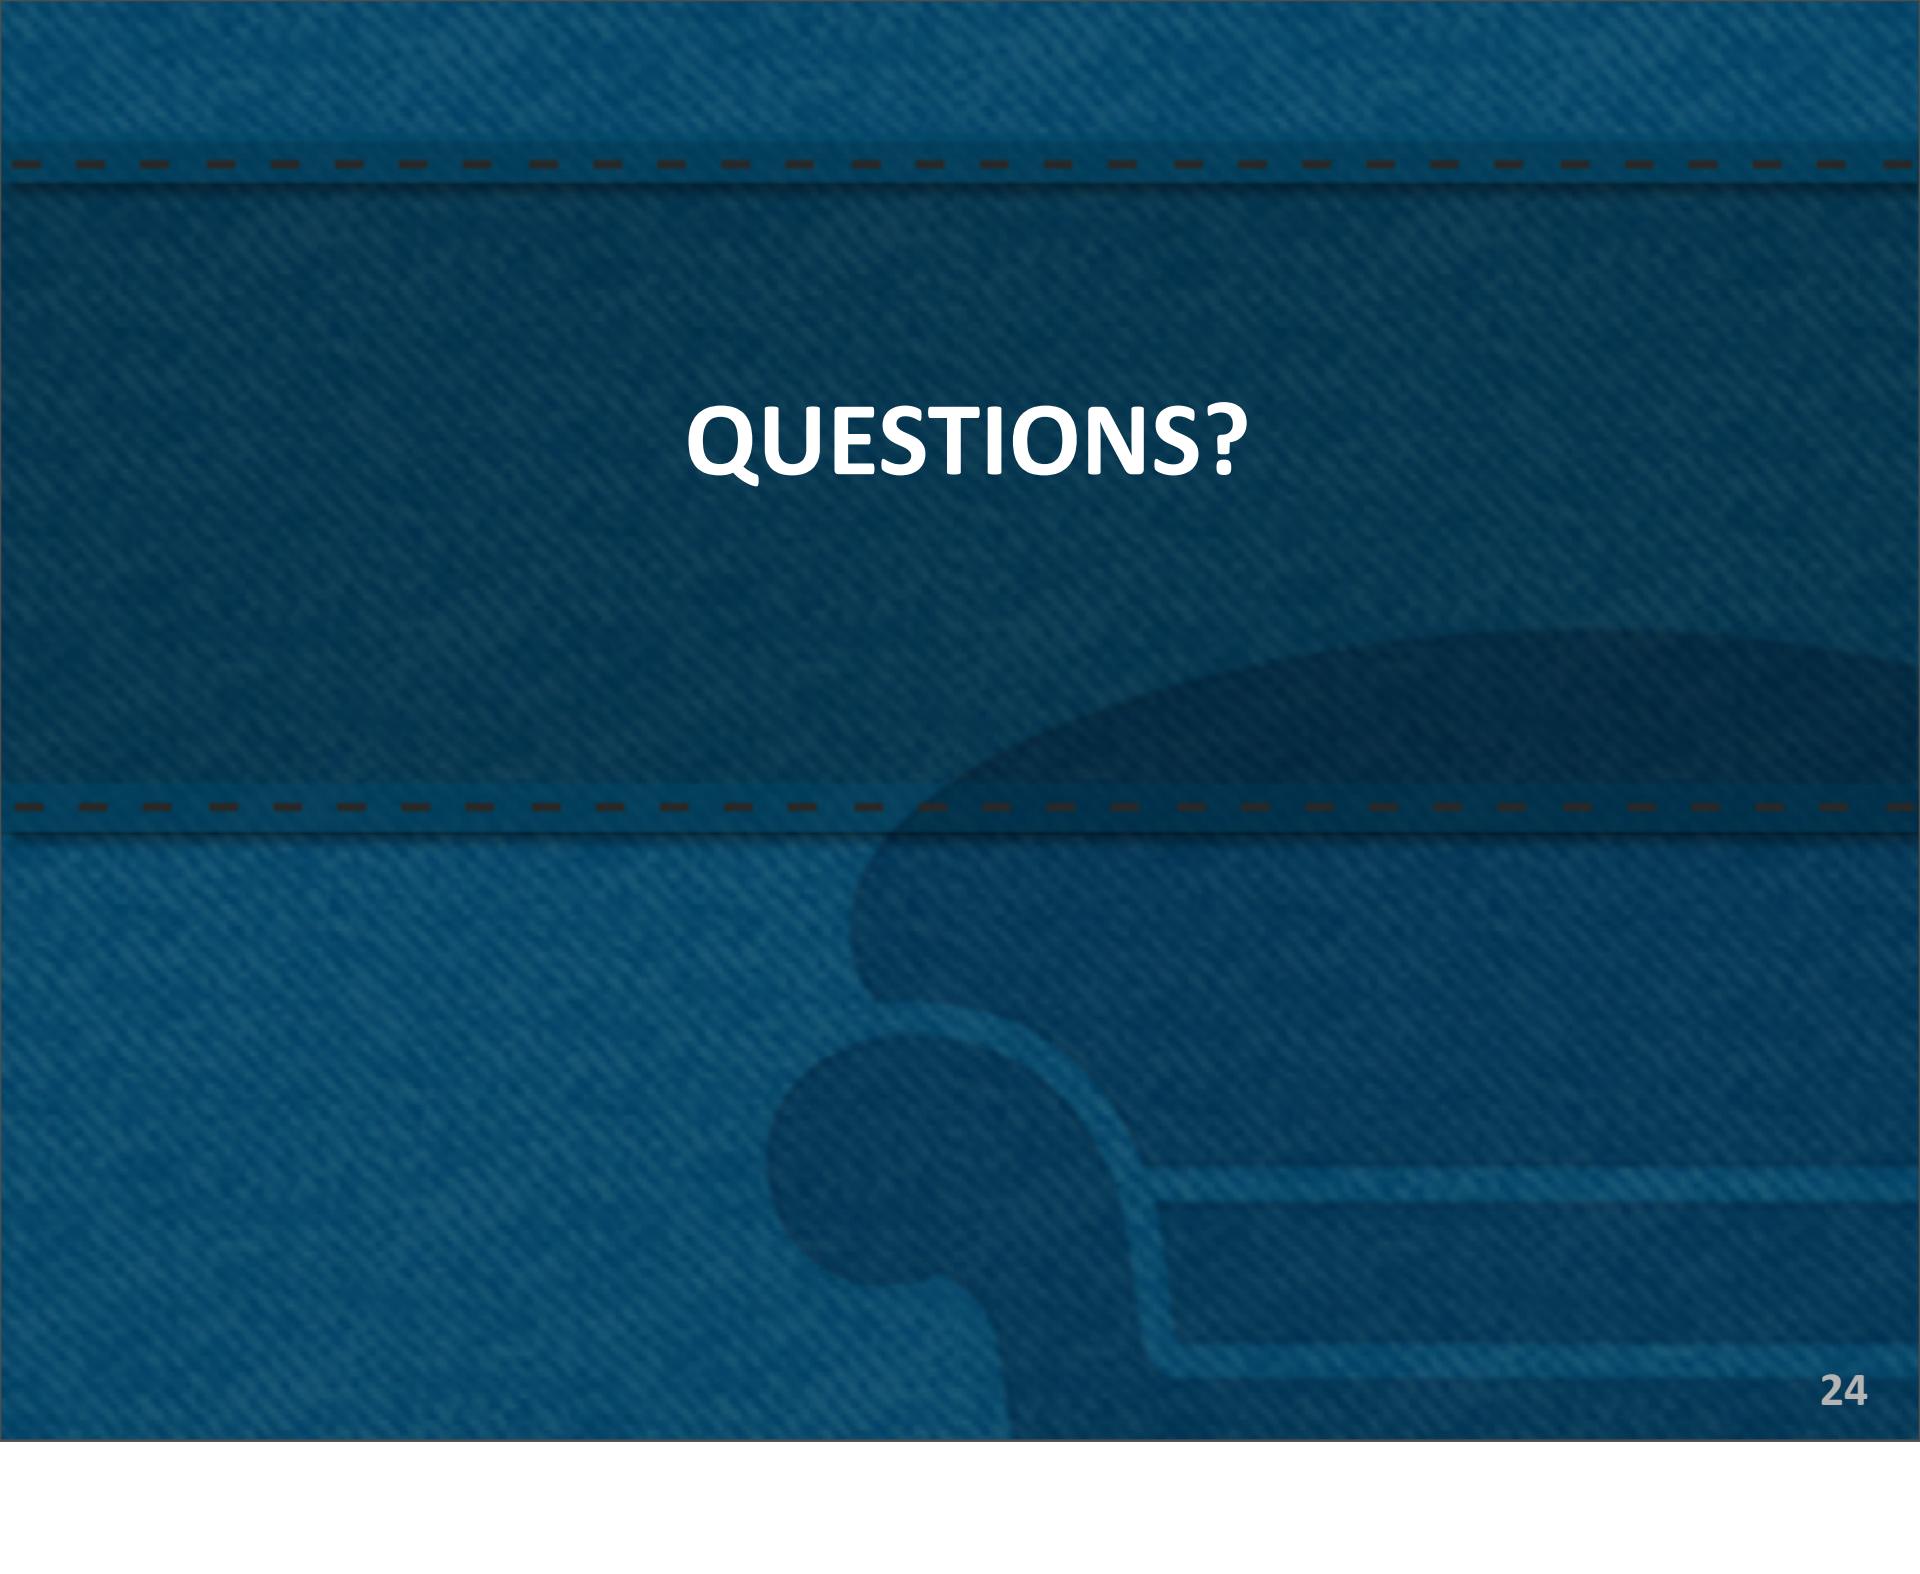
```
<?php  
$cb->set("a", 1); // "a" is 1  
$cb->increment("a"); // "a" is 2  
$cb->increment("a", 2); // "a" is 4  
$cb->decrement("a", 4); // "a" is 0
```

Arithmetic 2

```
<?php
// increment "b" by 3, create "b"
// if it doesn't exist in which
// case, initialise it with 2, set
// no expiry

// increment($key, $offset,
//           $create, $expiry, $initial)

$cb->increment("b", 3, true, 0, 2);
```



QUESTIONs?

THANK YOU!

Get Couchbase Server 2.0 at
<http://www.couchbase.com/downloads>

Give us feedback at:
<http://www.couchbase.com/forums>